

Large Language Models and Foundation Models for Time Series: A Survey and Outlook

Huirui Li

Abstract—Time series analysis underpins mission-critical decision systems across modern society, including industrial Internet of Things (IoT), computational finance, energy grid stabilization, meteorological forecasting, and clinical healthcare. Over the past five years, the field has undergone a profound paradigm shift from narrow, task-specific deep neural networks (such as specialized LSTMs and convolutional architectures) toward scalable foundation models. This evolution is bifurcated into two dominant architectural philosophies: (1) Native Time Series Foundation Models (TSFMs), which are trained ab initio on massive cross-domain time-series corpora comprising tens to hundreds of billions of observations (e.g., Chronos, TimesFM, MOIRAI, MOMENT, Lag-Llama, TTM, Timer, Sundial, Time-MoE, and Toto); and (2) Repurposed Large Language Models (LLM4TS), which repurpose frozen or parameter-efficiently adapted language transformers to temporal sequences through cross-modal token reprogramming, decimal prompting, or multimodal instruction tuning (e.g., Time-LLM, GPT4TS/OFA, LLMTime, TEMPO, and S2IP-LLM). In this survey, conducted strictly under the PRISMA 2020 systematic review protocol, we provide the first unified, comprehensive taxonomy and comparative synthesis spanning 2021 through 2026. We formalize the mathematical preliminaries of temporal representation, deconstruct core tokenization mechanics (patching vs. scalar quantization vs. lag vectors), analyze probabilistic formulation heads, examine empirical scaling behaviors and zero-shot transfer capabilities, and critically evaluate the central debate: Are billion-parameter language backbones truly necessary, or do compact specialized foundation models provide superior accuracy and operational efficiency? Finally, we conduct an audit of standardized benchmarks (e.g., GIFT-Eval), uncover evaluation pitfalls such as benchmark contamination and temporal leakage, and delineate open research directions toward multi-modal temporal intelligence.

Index Terms—Time series analysis, foundation models, large language models, deep learning, zero-shot forecasting, tokenization, scaling laws, PRISMA systematic review.



1 Introduction

Temporal data, systematically recorded as univariate or multivariate time series, serves as the fundamental perceptual substrate for understanding dynamic physical, economic, and biological systems. Across modern industries, time series forecasting, anomaly detection, imputation, and classification drive mission-critical workflows: stabilizing renewable energy power grids, optimizing multi-echelon retail supply chains, executing high-frequency financial risk arbitrage, predicting clinical sepsis in intensive care units, and monitoring aerospace telemetry [1].

Historically, time series modeling evolved through three distinct algorithmic epochs:

- 1) Statistical and Econometric Formulation (1970s–2010s): Dominated by classical stochastic processes, such as Auto-Regressive Integrated Moving Average (ARIMA), Exponential Smoothing (ETS), Vector Autoregression (VAR), and Gaussian Processes. While mathematically rigorous and interpretable, these methods rely on linear assumptions and fail to capture complex nonlinear dependencies across cross-domain series.
- 2) Task-Specific Deep Neural Networks (2015–2022): Catalyzed by recurrent architectures (LSTMs, GRUs), temporal convolutional networks (TCNs), and specialized

Transformer variants (Informer, Autoformer, FEDformer, and PatchTST [2]). Although these architectures captured long-range nonlinear temporal correlations, each model was trained in isolation on a single, isolated dataset (e.g., ETT, Weather, Electricity), exhibiting virtually zero out-of-distribution generalization.

- 3) The Foundation Model Epoch (2022–Present): Propelled by the extraordinary success of pre-trained foundation models in Natural Language Processing (e.g., GPT-4, Llama) and Computer Vision, the temporal machine learning community has initiated a paradigm shift toward universal, cross-domain foundation models [3]–[5]. Rather than training narrow models from scratch for every new dataset, practitioners now seek models capable of strong zero-shot and few-shot transfer across unseen domains, frequencies, and channels.

1.1 Unique Challenges of Temporal Signals

Transferring the foundation model paradigm to time series is non-trivial due to stark physical and mathematical differences between temporal signals and natural language:

- **Absence of a Discrete Natural Vocabulary:** Unlike natural text, which is composed of finite, discrete lexical tokens with stable semantic meanings, continuous time series $\mathbf{x} \in \mathbb{R}^T$ consist of real-valued scalar measurements bounded only by sensor physics.
- **Extreme Heterogeneity in Sampling Rates and Dynamics:** Real-world time series exhibit widely divergent temporal

• H. Li is the sole author and maintainer of this living survey repository. Contact: 58367737+lihuirui@users.noreply.github.com. Project repository: <https://github.com/lihuirui/awesome-llm-time-series-foundation-models-survey>.

scales—ranging from milliseconds in financial order books, to hourly power consumption, to daily retail demand, to quarterly economic indicators. A model must reconcile varying seasonalities without destructive aliasing.

- **Variable Dimensionality and Missing Values:** Time series data dynamically switch between univariate streams and arbitrary multivariate configurations ($C \geq 1$), frequently accompanied by irregular observation intervals, sensor failure dropouts, and structural trend shifts.

1.2 The Bifurcated Landscape: Native TSFMs vs. LLM4TS

To address these challenges, the research community has split into two competing architectural paradigms:

- 1) **Native Time Series Foundation Models (TSFMs):** Architectures designed and trained from scratch on massive multi-domain temporal corpora comprising tens to hundreds of billions of observations. Prominent exemplars include Amazon Chronos [3], [8], Google TimesFM [4], [9], Salesforce MOIRAI [5], [10], [11], CMU MOMENT [12], IBM TTM [13], Morgan Stanley/Mila Lag-Llama [14], and the THUML Timer family [15]–[18].
- 2) **Repurposed Large Language Models (LLM4TS):** Approaches that leverage the pre-trained weights, attention layers, and sequence modeling capabilities of multi-billion parameter language models (such as Llama, GPT-2, or GPT-4). Techniques include cross-modal token reprogramming (e.g., Time-LLM [19], GPT4TS/OFA [20], TEMPO [21]), direct numeric prompting with ASCII decimal strings (e.g., LLMTime [6]), and multimodal question-answering agents (e.g., Time-MQA [22], TimeOmni-1 [23]).

1.3 Key Contributions of This Survey

While several initial surveys have examined aspects of deep learning for time series, the rapid surge of foundation models from 2024 to 2026 demands a rigorous, systematic, and up-to-date synthesis. As highlighted in Table 1, this work delivers the following contributions:

- 1) **PRISMA 2020 Systematic Review Protocol:** We execute a strictly reproducible review protocol across scholarly APIs, logging 83 initial records, screening against standardized criteria, and synthesizing 41 core benchmark and foundation models.
- 2) **Rigorous Unified Taxonomy:** We establish an orthogonal five-dimensional taxonomy deconstructing models along: (i) architectural backbone, (ii) temporal tokenization, (iii) prediction head and uncertainty formulation, (iv) adaptation mechanics, and (v) operational task scope.
- 3) **Deep Comparative Synthesis of Native TSFMs vs. LLM4TS:** We analyze the theoretical and empirical tensions between compact native models (1M–700M parameters) and multi-billion parameter language adaptations (7B–70B parameters), examining computational efficiency, inference latency, and inductive bias alignment.
- 4) **Benchmark Audit, Data Contamination, and Scaling Laws:** We survey modern evaluation suites (e.g., GIFT-Eval [24]), expose systemic data leakage risks [25], analyze empirical power-law scaling in temporal models, and assess probabilistic calibration reliability [26].

- 5) **Open Research Agenda:** We articulate key open challenges, including irregular sampling equivariance, continuous flow matching, multimodal sensor fusion, and edge deployment.

1.4 Structure of the Survey

The remainder of this survey is organized as follows: Section 2 formalizes mathematical preliminaries and foundational objectives. Section 3 presents our comprehensive taxonomy and comparative design table. Section 4 deconstructs Native TSFMs. Section 5 investigates Repurposed LLM methods. Section 6 evaluates benchmark suites, scaling laws, and critical debates. Section 7 delineates open challenges, and Section 8 concludes the paper.

2 Background and Preliminaries

In this section, we establish the formal mathematical definitions, notation, and objective functions that govern modern temporal foundation models and large language model adaptations.

2.1 Formal Problem Definitions

2.1.1 Temporal Sequences

A univariate time series of length T is defined as an ordered sequence of real-valued scalar observations:

$$\mathbf{x} = [x_1, x_2, \dots, x_T]^\top \in \mathbb{R}^T, \quad (1)$$

where $x_t \in \mathbb{R}$ represents the measurement recorded at discrete timestamp $t \in \{1, \dots, T\}$.

A multivariate time series comprises C co-evolving channels or variates observed synchronously across T timestamps:

$$\mathbf{X} = [\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(C)}]^\top \in \mathbb{R}^{C \times T}, \quad (2)$$

where $\mathbf{X}_{:,t} \in \mathbb{R}^C$ is the vector of all variate measurements at timestamp t .

2.1.2 Channel Modeling Paradigms

When processing multivariate series $\mathbf{X}$, architectures adopt one of two primary structural paradigms:

- **Channel-Independence (CI):** Introduced prominently in PatchTST [2] and Timer [15], CI decomposes the multivariate series into C distinct univariate sequences:

$$\mathbf{X} \mapsto \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(C)}\}, \quad (3)$$

which are fed into the shared model backbone independently with tied parameters. CI eliminates the risk of overfitting spurious cross-channel correlations and natively handles variable channel counts C at inference.

- **Channel-Mixing (CM):** Models such as MOIRAI [5] and TTM [13] employ full cross-variate attention or specialized channel-mixing MLPs:

$$\mathbf{H} = \text{Attn}_{\text{channel}}(\mathbf{X}) \in \mathbb{R}^{C \times T \times D}, \quad (4)$$

explicitly capturing physical inter-sensor dependencies at the cost of higher computational complexity.

2.2 Tokenization and Input Representation

Because continuous temporal values lack an inherent lexicon, models deploy specialized mapping operators $\mathcal{T} : \mathbb{R}^* \rightarrow \mathbb{R}^D$ to construct transformer-compatible token embeddings.

TABLE 1
Comparison of This Survey with Related Surveys on Large Temporal Models.

Survey Reference	Year	PRISMA	Native TSFMs	Repurposed LLMs	Critiques & Audits	Primary Focus & Distinguishing Scope
Wen et al. [1]	2023	✗	✗	✗	✗	Survey of specialized Transformer architectures on single datasets.
Jin et al. [1]	2024	✗	Partial	✓	Partial	Broad survey covering both general time series and spatio-temporal graphs.
Gruver et al. [6]	2023	✗	✗	✓	✗	Empirical investigation of prompting LLMs as zero-shot forecasters.
Fons et al. [7]	2024	✗	✗	✓	✓	Empirical critique testing whether language backbones add genuine inductive value.
This Survey	2026	✓	✓	✓	✓	Comprehensive PRISMA systematic review (2021–2026); unifying Native TSFMs, Repurposed LLMs, tokenization trade-offs, scaling laws, and leakage audits.

2.2.1 Subseries Patching

Rather than embedding individual scalar time-steps (which leads to quadratic attention overhead $\mathcal{O}(T^2)$ and severe information sparsity), the majority of modern TSFMs (e.g., TimesFM [4], Timer [15], MOMENT [12]) partition a sequence $\mathbf{x}$ into N contiguous subseries patches of length P with stride S :

$$N = \left\lfloor \frac{T-P}{S} \right\rfloor + 1. \quad (5)$$

The i -th patch $\mathbf{p}_i = [x_{(i-1)S+1}, \dots, x_{(i-1)S+P}]^\top \in \mathbb{R}^P$ is mapped to a D -dimensional latent vector $\mathbf{e}_i \in \mathbb{R}^D$ via an affine projection:

$$\mathbf{e}_i = \mathbf{p}_i \mathbf{W}_p + \mathbf{b}_p + \mathbf{e}_{\text{pos},i}, \quad (6)$$

where $\mathbf{W}_p \in \mathbb{R}^{P \times D}$, $\mathbf{b}_p \in \mathbb{R}^D$, and $\mathbf{e}_{\text{pos},i} \in \mathbb{R}^D$ denotes a learnable or sinusoidal positional encoding. Patching preserves local temporal shape dynamics and compresses sequence length from T to $N \ll T$.

2.2.2 Scalar Quantization and Categorical Binning

In contrast to continuous linear patching, Amazon Chronos [3] introduces a pure language-analog tokenization. Given an input sequence $\mathbf{x}$, it computes instance-level scale statistics (μ, σ) via mean scaling:

$$\tilde{x}_t = \frac{x_t - \mu}{\sigma + \epsilon}, \quad \mu = \frac{1}{T} \sum_{t=1}^T x_t, \quad \sigma = \frac{1}{T} \sum_{t=1}^T |x_t - \mu|. \quad (7)$$

The normalized scalars $\tilde{x}_t$ are discretized into K uniform bins over $[-B, B]$:

$$q_t = \text{clip} \left(\left\lfloor \frac{\tilde{x}_t - (-B)}{2B/K} \right\rfloor, 0, K-1 \right), \quad (8)$$

where $q_t \in \{0, 1, \dots, K-1\}$ acts as a discrete categorical token id, enabling standard cross-entropy language modeling.

2.3 Downstream Task Formulations

2.3.1 Zero-Shot and Supervised Forecasting

Given an historical context window $\mathbf{X}_{1:T} \in \mathbb{R}^{C \times T}$, the forecasting objective is to predict the future sequence of length H :

$$\hat{\mathbf{Y}} = \hat{\mathbf{X}}_{T+1:T+H} \in \mathbb{R}^{C \times H}. \quad (9)$$

In probabilistic forecasting, the model estimates the conditional predictive distribution:

$$p(\mathbf{Y} \mid \mathbf{X}_{1:T}) = \prod_{h=1}^H p(y_{T+h} \mid \mathbf{X}_{1:T}, y_{T+1:T+h-1}). \quad (10)$$

This distribution is parameterized either parametrically (e.g., Student's t parameters in Lag-Llama [14] or Gaussian mixtures in MOIRAI [5]) or non-parametrically through categorical bin logits (Chronos [3]).

2.4 Pre-training Objectives

2.4.1 Autoregressive Generation (Next-Patch / Next-Token)

Decoder-only foundation models (TimesFM [4], Timer [15], Chronos [3]) optimize an autoregressive likelihood:

$$\mathcal{L}_{\text{AR}}(\Theta) = - \sum_{i=1}^N \log p_{\Theta}(\mathbf{p}_{i+1} \mid \mathbf{p}_1, \dots, \mathbf{p}_i). \quad (11)$$

For continuous regression heads (Timer, TimesFM), this is computed using Mean Squared Error (MSE) or Mean Absolute Error (MAE):

$$\mathcal{L}_{\text{reg}}(\Theta) = \frac{1}{N} \sum_{i=1}^N \|\mathbf{p}_{i+1} - \hat{\mathbf{p}}_{i+1}\|_2^2. \quad (12)$$

2.4.2 Masked Autoencoding (MAE)

Encoder-based foundation models (MOMENT [12], PatchTST [2]) apply a binary masking vector $\mathbf{m} \in \{0, 1\}^N$ over patches, optimizing patch reconstruction over the masked set $\mathcal{M} = \{i \mid m_i = 1\}$:

$$\mathcal{L}_{\text{Mask}}(\Theta) = \frac{1}{|\mathcal{M}|} \sum_{i \in \mathcal{M}} \|\mathbf{p}_i - f_{\Theta}(\mathbf{X}_{\setminus \mathcal{M}})\|_2^2. \quad (13)$$

Masked autoencoding yields bidirectional temporal representations particularly suited for downstream classification, anomaly detection, and imputation.

3 Taxonomy and Landscape Overview

To organize the rapid proliferation of temporal foundation models and large language model adaptations, we present an orthogonal multi-dimensional taxonomy. As illustrated in Figure 1, the landscape is structured along five fundamental axes: (1) Core Paradigm, (2) Architectural Backbone, (3) Temporal Tokenization, (4) Prediction Head & Uncertainty, and (5) Operational Scope & Tasks.

3.1 Taxonomy Dimensions

3.1.1 Dimension 1: Core Paradigm

- **Native Time Series Foundation Models (Native TSFMs):** Models built from scratch and trained directly on large-scale collections of empirical or synthetic time-series data. Weights and latent dimensions are optimized specifically for temporal signals without textual preconceptions (e.g., Chronos [3], TimesFM [4], MOIRAI [5], ForecastPFN [27], TiRex [28], FLAME [29], t_0 [30]).
- **Repurposed Large Language Models (LLM4TS):** Models that adopt pre-existing weights from language or vision transformers (e.g., Llama-2/3, GPT-2, CLIP), bridging continuous temporal modalities into text or visual embedding space via reprogramming layers, prompt prefixes, or cross-attention adapters (e.g., Time-LLM [19], GPT4TS [20], TEST [31], CALF [32], UniTime [33], VisionTS [34], Time-VLM [35]).
- **Evaluation Frameworks and Benchmark Audits:** Standardized benchmarking suites and diagnostic harnesses that audit generalization, cross-domain transfer, and benchmark contamination (e.g., GIFT-Eval [24], SciTS [36], Insight Miner [37], Time-MMD [38]).

3.1.2 Dimension 2: Architectural Backbone

- **Decoder-Only Autoregressive Transformers:** Optimized for sequential next-token or next-patch generation (e.g., TimesFM [4], Timer [15], Chronos [3], Sundial [17], t_0 [30]).
- **Masked Encoder Transformers:** Optimized for bidirectional temporal representations via masked autoencoding (e.g., MOMENT [12], PatchTST [2], VisionTS [34]).
- **Encoder-Decoder Transformers:** Decoupling context representation from horizon generation (e.g., MOIRAI [5], TTM [13], UniTS [39]).
- **Sparse Mixture-of-Experts (MoE):** Dynamically routing tokens to specialized feed-forward experts to scale model capacity while maintaining constant inference FLOPs (e.g., Time-MoE [40], Moirai-MoE [10], Timer-S1 [18]).
- **Continuous Flow and Legendre Memory Networks:** Employing continuous neural ODEs or orthogonal polynomial

bases to preserve multi-scale memory (e.g., FLAME [29], FlowState [41]).

3.1.3 Dimension 3: Temporal Tokenization

- **Uniform Subseries Patching:** Continuous linear projection of non-overlapping or strided time-step windows into dense embedding space (e.g., PatchTST [2], Timer [15]).
- **Scalar Discretization / Binning:** Quantizing scaled continuous observations into finite categorical tokens, allowing standard NLP cross-entropy training (Chronos [3]).
- **Decimal String Tokenization:** Casting continuous numbers into ASCII strings (e.g., "12.4"), processed directly by off-the-shelf text tokenizers (PromptCast [42], LLMTime [6]).
- **Multi-Resolution and Multiscale Mixing:** Dynamically decomposing and mixing patches across frequencies to resolve multi-scale seasonalities (TimeMixer [43], TimeMixer++ [44], TTM [13]).
- **Visual Rendering Tokenization:** Rendering 1D numerical trajectories into 2D pixel line plots for visual encoders (VisionTS [34], Time-VLM [35]).

3.1.4 Dimension 4: Prediction Head and Uncertainty

- **Categorical Probability Distribution:** Non-parametric probability vectors over quantized bins (Chronos [3]).
- **Parametric Mixture Heads:** Estimating parameters (μ, σ, ν) of explicit probability densities, such as Student's t or Gaussian mixtures (Lag-Llama [14], MOIRAI [5]).
- **Deterministic Mean Squared Regression:** Direct linear or MLP projection minimizing mean squared or absolute error.
- **Continuous Flow Matching and Diffusion:** Formulating generation as an equivariant continuous-time velocity field (FlowState [41], FLAME [29]).
- **Quantile Regression Heads:** Directly outputting discrete quantile trajectories to provide non-parametric prediction intervals (t_0 [30], TimesFM [4], Toto 2.0 [45]).

4 Native Time Series Foundation Models

Native Time Series Foundation Models (TSFMs) represent the purist approach to temporal foundation modeling: architectures designed from first principles and trained directly on massive corpora of continuous time-series data without linguistic baggage. In this section, we examine the primary architectural branches, data curation pipelines, and empirical scaling mechanisms.

4.1 Autoregressive Decoder-Only Architectures

Autoregressive models frame time-series modeling as sequential generation, predicting subsequent time-steps or patches conditioned on historical context.

4.1.1 The Chronos Family: Discretization as Linguistic Analogy

Amazon's Chronos [3] represents one of the most radical departures from conventional continuous time-series modeling. Rather than predicting real numbers via regression losses, Chronos asks: Can time series be modeled identically to language tokens? By applying instance-level mean scaling followed by uniform quantization into $K = 4096$ bins, continuous values are transformed into integer token indices.

Taxonomy of Foundation Models and Large Language Models for Time Series




Fig. 1. Comprehensive taxonomy of foundation models and large language models for time series, illustrating core paradigms, architectural backbones, tokenization mechanics, and representative models.

TABLE 2
Comprehensive Comparison of Representative Time Series Foundation Models and LLM Methods.

Model	Group	Venue/Yr	Paradigm	Backbone	Tokenization	Params	Pretraining Corpus	Uncertainty	Code / Status
Chronos [3]	Amazon	TMLR'24	Native TSFM	Decoder	Quantized Bins	20M–710M	TSMix (84B obs)	Categorical	amazon/chronos
TimesFM [4]	Google	ICML'24	Native TSFM	Decoder	Patch ($P = 32$)	200M	Real+Synth (100B)	Quantile	google/timesfm
MOIRAI [5]	Salesforce	ICML'24	Native TSFM	Enc-Dec	Multi-patch (8–128)	14M–311M	LOTSa (27B obs)	Mixture Param	salesforce/uni2ts
MOMENT [12]	CMU	ICML'24	Native TSFM	Encoder	Patch ($P = 512$)	385M	TS-Pile (1B obs)	Recon Head	cmu/moment
Lag-Llama [14]	Mila / MS	ICML'24	Native TSFM	Decoder	Lag Vectors	2.4M	Monash (7.9K ser)	Student's t	mila/lag-llama
TTM [13]	IBM	NeurIPS'24	Native TSFM	Enc-Dec	Adaptive Patch	1M–8M	Multi-domain Pub	Deterministic	ibm/granite-tsfm
Timer [15]	THUML	ICML'24	Native TSFM	Decoder	Subseries Patch	84M	UTSD (1B pts)	Next-Patch	thuml/Timer
Timer-XL [16]	THUML	NeurIPS'24	Native TSFM	Decoder	Long-Ctx Patch	84M	UTSD-2	Next-Patch	thuml/Timer-XL
Sundial [17]	THUML	arXiv'25	Native TSFM	Decoder	Multi-rate Patch	1.5B	UTSD-3 (>10B)	Continuous	thuml/Sundial
Time-MoE [40]	Ming Jin	ICLR'25	Native TSFM	MoE	Multi-res Patch	2.4B	Time-300B (300B)	MoE Contin	Time-MoE
Timer-S1 [18]	THUML	arXiv'26	Native TSFM	MoE	Serial Patch	8.3B	UTSD-3	MoE Regress	thuml/Timer-S1
ForecastPFN [27]	Abacus/CMU	NeurIPS'23	Native TSFM	Transformer	In-Context PFN	5M	Synthetic Prior	Quantile	abacus/ForecastPFN
t_0 [30]	Invenia/ETH	arXiv'26	Native TSFM	Axial Attn	Context Patch	102M–256M	Pub+Synth	Quantile	Proprietary
Toto 2.0 [45]	Datadog	arXiv'26	Native TSFM	Decoder	Wavelet Tokens	1.2B	Telemetry (1.5T)	Quantile	Proprietary
Time-LLM [19]	Ming Jin	ICLR'24	LLM4TS	Llama-7B	Reprogrammed	7B	Pretrained LLM	Deterministic	KimMeen/Time-LLM
GPT4TS [20]	DAMO	NeurIPS'23	LLM4TS	GPT-2	Cross-Modal	Unreported	Pretrained GPT-2	Deterministic	DAMO/GPT4TS
TEST [31]	PKU/Alibaba	NeurIPS'24	LLM4TS	Llama/BERT	Text Prototype	7B	Pretrained LLM	Deterministic	Proprietary
CALF [32]	THU/DAMO	KDD'24	LLM4TS	Llama-7B	Distilled Patch	7B	Pretrained LLM	Deterministic	Hank0626/CALF
LLMTime [6]	NYU	NeurIPS'23	LLM4TS	GPT-4/Llama	Decimal ASCII	70B+	Pretrained LLM	Sampling Logit	ngruver/llmtime
TEMPO [21]	Ming Jin	ICLR'24	LLM4TS	GPT-2	Decomposed	Unreported	Pretrained GPT-2	Decomposed	caocao/TEMPO
AutoTimes [46]	THUML	NeurIPS'24	LLM4TS	Llama-2-7B	Point-wise	7B	Pretrained LLM	Autoregress	thuml/AutoTimes
Time-MQA [22]	Ming Jin	ACL'25	LLM4TS	Llama-3-8B	Patch + Prompt	8B	Multi-domain MQA	Text/Value	Proprietary

Chronos trains standard T5 architectures (ranging from 20M to 710M parameters) solely with the cross-entropy objective. The predictive distribution over future horizons is obtained by autoregressively sampling categorical bin logits and dequantizing them back into real space. This approach has three key advantages:

- 1) It sidesteps parametric distribution assumptions (e.g., Gaussianity), naturally expressing multi-modal, skewed, or heavy-tailed predictive distributions.
- 2) It enables direct reuse of optimized language transformer implementations (such as FlashAttention and KV-caching).
- 3) In Chronos-2 [8], the architecture is extended to multi-

variate series via grouped channel tokenization and hybrid continuous residual heads.

4.1.2 Google TimesFM: Multi-Patch Output Generation

Unlike Chronos's scalar quantization, Google's TimesFM [4] operates directly on continuous numerical values using continuous patching. TimesFM employs an input patch size of $P = 32$ and an output patch size of $H = 128$. By predicting an entire multi-step block in a single forward step, TimesFM drastically reduces autoregressive unrolling steps, curbing error accumulation and inference latency. TimesFM is trained on a massive corpus of 100B time-series data points collected from Google Trends, Wikipedia pageviews, and synthetic

ARMA processes. Recent extensions examine in-context fine-tuning [9], demonstrating that supplying short conditioning examples within the context window improves zero-shot adaptation. Subsequently, TimesFM-3 scales the architecture to 330M parameters pre-trained over 1T time points, introducing native multivariate forecasting via stacked variate attention and iterative reversible instance normalization (RevIN). Demonstrating the industrial viability of large-scale foundation forecasters, Tacconelli [47] introduces Cadence, pairing TimesFM-3 with an adaptive arithmetic coder to execute error-bounded lossy compression on enterprise telemetry, guaranteeing user-defined maximum deviation bounds while achieving state-of-the-art bit-rate reductions.

4.1.3 The THUML Timer Family and Output Scaling

The THUML research group (led by Mingsheng Long at Tsinghua University) has systematically advanced generative time series pretraining:

- Timer [15]: Conceptualizes a single-series sequential next-patch regression paradigm. Timer is pretrained on the Unified Time Series Dataset (UTSD), comprising over 1B observations across diverse domains. Timer establishes that channel-independent single-series formatting enables universal transfer across arbitrary multivariate configurations.
- Timer-XL [16]: Introduces long-context transformer modeling for temporal signals, scaling context lengths from standard 512 patches to extreme contexts, enabling the model to capture multi-year periodicities without context truncation.
- Sundial [17]: Scales the model to 1.5B parameters pretrained on UTSD-3 (> 10B observations), introducing multi-rate adaptive patching to handle heterogeneous sampling frequencies seamlessly.
- Timer-S1 [18]: Reaches 8.3B parameters via sparse Mixture-of-Experts routing, demonstrating the viability of serial parameter scaling in native temporal models.
- YingLong [48]: Explores test-time output scaling in pre-trained forecasting models by introducing delayed chain-of-thought (CoT) tokens. Rather than predicting future values in a single forward pass, YingLong autoregressively generates intermediate trend and regime change tokens before decoding target values, demonstrating that computational output scaling laws apply to temporal reasoning.

4.1.4 Context-Conditioned and Dual-Horizon Autoregressive Models

Standard autoregressive models often struggle with the structural tension between resolving high-frequency short-term fluctuations and maintaining long-range macro trajectories. Recent 2025–2026 architectures address this via explicit context-conditioning:

- TiRex [28]: Auer et al. (JKU Linz / ELLIS) propose an asymmetric dual-horizon in-context architecture. TiRex isolates local fine-grained patches from global horizon summaries, using memory compaction to preserve multi-scale context without quadratic computational blowup.
- t_0 Family [30]: Meyer et al. introduce t_0 , releasing open-weights checkpoints t_0 -alpha (102M) and t_0 -beta (256M). Departing from channel-independent univariate approaches, t_0 explicitly conditions predictions on historical

targets, past dynamic covariates, and known-future covariates (e.g., promotional calendars, weather forecasts). By alternating axial self-attention along the temporal axis and cross-variate attention across feature channels, t_0 attains aggregate CRPS of 0.4738 and MASE of 0.6865 on GIFT-Eval, standing within 4% of the top zero-shot forecaster.

- Toto Family [45], [49]: Datadog introduced Toto (151M parameters) [49], trained on 1T cloud infrastructure and application metrics to deliver joint zero-shot quantile forecasting and out-of-distribution anomaly scoring. Toto 2.0 [45] subsequently scales from 4M to 2.5B parameters using Maximal Update Parametrization (μ P), achieving state-of-the-art results on GIFT-Eval and the contamination-resistant TIME benchmark.
- Tabby Open Recipe [50]: Addressing the pervasive opacity of proprietary pretraining mixtures, Xie et al. introduce Tabby (120M parameters), releasing an open-source, long-context probabilistic foundation model accompanied by a fully reproducible data curation pipeline and an open 150B-token pretraining corpus.

4.1.5 Lag-Llama: Probabilistic Modeling with Lag Embeddings

Developed by Morgan Stanley and Mila, Lag-Llama [14] adapts the Llama decoder backbone by substituting word embeddings with lag feature vectors. Specifically, it computes lags corresponding to standard calendar frequencies (e.g., quarter-hour, day, week, year). Lag-Llama outputs the parameters of a Student’s t distribution, achieving robust probabilistic calibration on out-of-distribution Monash datasets.

4.2 Prior-Data Fitted Networks (PFNs) and In-Context Meta-Learning

A radically distinct paradigm in native foundation modeling is the Prior-Data Fitted Network (PFN) and in-context learning framework, which foregoes empirical training corpora or iterative gradient steps entirely in favor of synthetically generated structural priors:

- ForecastPFN [27]: Dooley et al. demonstrate that a transformer trained exclusively on synthetic time series generated from a Bayesian mixture of Gaussian processes, ARMA models, and dynamical equations can approximate posterior predictive distributions in a single forward pass. ForecastPFN requires zero gradient steps or hyperparameter tuning at inference time, achieving competitive zero-shot accuracy across diverse real-world datasets.
- TabPFN-TS [51]: Extending tabular foundation models (TabPFN-v2) to the temporal domain, Müller et al. propose TabPFN-TS (11M parameters). By encoding time series as tabular regression tasks conditioned on causal lag features and exogenous metadata, TabPFN-TS achieves state-of-the-art accuracy on covariate-informed forecasting across GIFT-Eval and fev-bench, proving that in-context synthetic priors offer extreme parameter efficiency.
- In-Context Time Series Predictor (ICTSP) [52]: Lu et al. (ICLR 2025) formalize in-context regression for continuous temporal sequences. By structuring historical context and query sequences within an interleaved prompt-query attention mechanism, ICTSP (42M parameters) executes zero-gradient meta-adaptation across disparate sampling

frequencies and noise regimes without updating network weights.

4.3 Masked Encoders, Multiscale Mixers, and Multi-Task Architectures

4.3.1 Salesforce MOIRAI: Any-Variate Cross-Attention

Salesforce’s MOIRAI [5] addresses two central limitations of previous models: rigid patch sizes and channel-independence. MOIRAI introduces Any-Variate Attention, flattening an arbitrary multivariate sequence $\mathbf{X} \in \mathbb{R}^{C \times T}$ into a sequence of variate-tagged patch tokens. Furthermore, it employs a multi-patch projection head capable of simultaneously processing patch lengths of 8, 16, 32, 64, and 128. MOIRAI is trained on the Large-scale Open Time Series Archive (LOTSAs), containing 27B observations across 9 distinct domains. The model outputs mixtures of parametric distributions, achieving top-tier ranking on zero-shot benchmarks. Moirai 2.0 [11] and Moirai-MoE [10] further extend this framework to sparse routing with 1.1B parameters.

4.3.2 CMU MOMENT and UniTS: Universal Multi-Task Representation

While most TSFMs focus exclusively on forecasting, CMU’s MOMENT [12] investigates foundation models for universal multi-task analysis (forecasting, classification, anomaly detection, and imputation). Based on a masked T5 encoder (385M parameters), MOMENT is pretrained on the Time-series Pile (1B observations) using masked patch reconstruction. Similarly, Harvard’s UniTS [39] unifies multi-task modeling through a single transformer backbone. By combining shared attention operators with task-specific prompt vectors and dynamic gating networks, UniTS demonstrates that cross-task representations mutually reinforce classification and forecasting generalization.

4.3.3 Lightweight Mixers and Exogenous Variable Modeling

Challenging the premise that foundation models require massive parameter scales:

- Tiny Time Mixers (TTM) [13]: IBM Research introduced TTM (1M–8M parameters), utilizing lightweight multi-level TSMixer blocks with adaptive resolution sampling. TTM out-executes models orders of magnitude larger while running inference in milliseconds on standard edge CPUs.
- TimeMixer and TimeMixer++ [43], [44]: Wang et al. formulate multiscale mixing architectures. TimeMixer decomposes signals into downsampled representations to disentangle multiscale seasonalities, while TimeMixer++ extends this principle into a general temporal pattern machine supporting cross-domain transfer.
- TimeXer [53]: Specifically targets exogenous variables by learning cross-attention maps between target end-points and exogenous features, avoiding noise propagation from uninformative covariates.
- Kairos [54]: Introduces predictability-stratified pretraining (PreSTS), filtering chaotic or non-deterministic temporal noise during self-supervised training to ensure foundation models focus representational capacity on predictable temporal dynamics.
- LightGTS [55]: Wang et al. (ICML 2025) challenge the assumption that cross-domain generalizability requires heavy

parameter footprints. LightGTS (5.8M parameters) introduces dual-path hierarchical patching and adaptive parameter allocation, matching the zero-shot accuracy of foundation models two orders of magnitude larger while running inference directly on resource-constrained edge microcontrollers.

4.4 Continuous Flow Matching and Polynomial Memory Networks

A vital 2025 frontier replaces discrete tokenization with continuous-time dynamical systems:

- FLAME [29]: Wu et al. introduce Flow-Enhanced Legendre Memory Models. FLAME compresses long continuous temporal histories onto orthogonal Legendre polynomial bases, avoiding memory degradation over extended sequences, while modeling predictive distributions via continuous normalizing flows.
- FlowState [41]: King et al. formulate time series forecasting via continuous-time ODEs and flow matching. FlowState achieves strict sampling-rate equivariance, allowing zero-shot inference across irregularly sampled or arbitrary continuous time-steps without interpolation artifacts. Despite its compact footprint (12M parameters), FlowState achieves top-tier CRPS on GIFT-Eval.
- EIDOS [56]: Zhou et al. (Ming Jin’s group, 2026) introduce latent-space predictive learning for temporal foundation models via a Joint-Embedding Predictive Architecture (JEPA). Rather than optimizing token-level mean-squared reconstruction—which forces the model to expend capacity memorizing stochastic high-frequency jitter—EIDOS predicts future dynamics within a learned latent metric space, yielding semantically structured representations that transfer robustly across unseen domains.

4.5 Mixture-of-Experts (MoE) and Scaling to Billions

A major architectural breakthrough between 2024 and 2026 is the deployment of Sparse Mixture-of-Experts (MoE) in time series. In Time-MoE [40] (Ming Jin’s group), the model scales to 2.4B parameters while only activating a subset of experts per patch token. Trained on Time-300B (300B time points), Time-MoE demonstrates clear empirical power-law scaling in validation loss, proving that scaling laws discovered in NLP hold equally for native temporal representations when sufficient data volume is supplied. Similarly, Moirai-MoE [10] (1.1B) and Timer-S1 [18] (8.3B) substantiate that sparse MoE routing allows foundation models to specialize on domain-specific dynamics (e.g., high-frequency finance vs. slow-moving climate trends) while keeping inference latency within practical bounds.

4.6 Parameter-Efficient Fine-Tuning vs. In-Context Adaptation

A central dilemma in deploying foundation models is choosing between zero-shot inference, in-context learning, parameter-efficient fine-tuning (PEFT), and full fine-tuning. Table 3 categorizes the core adaptation paradigms across efficiency, computational overhead, and empirical generalization trade-offs.

Empirical synthesis reveals two critical principles governing downstream adaptation:

TABLE 3
Systematic Comparison of Downstream Adaptation Paradigms for Time Series Foundation Models.

Adaptation Paradigm	Exemplar Models	Trainable Params	Gradient Steps	Memory Overhead	Data Regime / Sample Efficiency	Generalization & Stability
Zero-Shot Transfer	Chronos-2 [8], TimesFM [4], t_0 [30], Toto 2.0 [45]	0%	0	Baseline (KV-Cache only)	Zero target samples required; immediate out-of-the-box deployment	Preserves broad cross-domain calibration; bounded by pretraining domain overlap
In-Context Meta-Learning	ForecastPFN [27], TabPFN-TS [51], ICTSP [52], TimesFM-IC [9]	0%	0	Low-Medium (Expanded context window)	Extreme few-shot (5-50 context sequences); zero fine-tuning FLOPs	Immune to catastrophic forgetting; limited by maximum attention context length
Parameter-Efficient (PEFT / LoRA / Prefix)	TTM [13], CALF [32], Time-LLM [19], OFA [20]	0.05%-2.0%	Few (50-500)	Low (< 10% activation memory)	Highly sample-efficient (5%-10% target dataset matches full training)	Mitigates representation collapse via low-rank constraints; retains core backbone weights
Full-Parameter Fine-Tuning	MOMENT [12], PatchTST [2], Timer [15]	100%	Many (1K-10K)	High (4× model size for optimizer states)	Requires substantial target history (100% dataset); prone to overfitting on small data	Achieves lowest in-domain MSE on large stationary datasets; suffers catastrophic forgetting

- 1) The In-Context Learning Frontier: Prior-Data Fitted Networks (TabPFN-TS, ForecastPFN) and meta-learning predictors (ICTSP) demonstrate that conditioning on prompt-target pairs inside the attention context matches 100-step gradient descent while executing in a single forward pass, providing an ideal mechanism for real-time edge streaming where backpropagation is computationally prohibited.
- 2) PEFT Prevents Semantic Collapse: As demonstrated in CALF [32] and TTM [13], tuning less than 1% of weights via low-rank matrices or adaptive resolution prefix blocks yields within 2% of full fine-tuning performance on target benchmarks while completely insulating the model against catastrophic representation drift.

4.7 Spatio-Temporal Foundation Models: Unifying 1D Streams and Graph Topologies

A pervasive frontier in foundation modeling is extending beyond isolated 1D univariate or channel-independent series toward structured spatio-temporal systems. In transportation networks, urban human mobility, power distribution grids, and planetary meteorological monitoring, continuous observation streams are inextricably situated upon complex underlying geometric topologies [57].

Formally, a spatio-temporal system is represented as a spatial topology $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathbf{A})$, where $\mathcal{V}$ denotes a set of $N = |\mathcal{V}|$ spatial sensor nodes or grid cells, $\mathcal{E}$ represents physical or learned relational edges, and $\mathbf{A} \in \mathbb{R}^{N \times N}$ is the spatial adjacency matrix (or spatial distance metric). Given historical observations $\mathbf{X}_{1:T} = [\mathbf{x}_1, \dots, \mathbf{x}_T] \in \mathbb{R}^{T \times N \times D_{\text{in}}}$, the spatio-temporal foundation objective requires forecasting future states $\mathbf{X}_{T+1:T+H} \in \mathbb{R}^{H \times N \times D_{\text{out}}}$ across both seen and entirely unseen spatial topologies $\mathcal{G}^*$:

$$P(\mathbf{X}_{T+1:T+H} \mid \mathbf{X}_{1:T}, \mathcal{G}) = \prod_{h=1}^H P(\mathbf{x}_{T+h} \mid \mathbf{X}_{1:T+h-1}, \mathbf{A}). \quad (14)$$

Traditional deep spatio-temporal architectures (e.g., STGCN, DCRNN) relied on fixed static spatial graph Laplacian operators and task-specific training, catastrophically failing when deployed zero-shot across new sensor networks or different urban geographies. Modern Spatio-Temporal Foundation Models (STFMs) circumvent this structural barrier through four primary paradigms:

- Dual Spatio-Temporal Self-Attention with Graph Biases (OpenCity): Li et al. [58] present OpenCity, an open-world

foundation model for urban traffic forecasting. OpenCity decomposes representations into decoupled spatial and temporal transformer layers. To eliminate reliance on rigid fixed adjacency matrices, OpenCity injects generalized graph wavelet positional encodings and relative spatial distance embeddings into multi-head self-attention. Pre-trained on multi-city traffic networks comprising over 50M observations, OpenCity demonstrates remarkable zero-shot transfer across previously unobserved metropolitan road networks.

- Prompt-Empowered Universal Masked Autoencoding (UniST): Yuan et al. [59] introduce UniST (Tsinghua FIB Lab), moving beyond scenario-specific architectures. UniST formulates spatio-temporal forecasting through unified generative masked pretraining. To accommodate heterogeneous sampling rates, variable sensor densities, and semantic domain shifts (e.g., traffic velocity vs. pedestrian flow vs. bike sharing), UniST injects knowledge-guided continuous prompt tokens into the transformer backbone, guiding the shared model to dynamically adapt spatial-temporal priors without altering backbone parameters.
- Spatio-Temporal Diffusion Transformers (UrbanDiT): Li et al. [60] establish UrbanDiT, scaling Diffusion Transformers (DiT) to open-world spatio-temporal systems. UrbanDiT tokenizes both irregular spatial graphs and regular spatial raster grids into unified spatio-temporal patches. By combining reverse-time score-based diffusion denoising with task prompt embeddings, UrbanDiT unifies diverse downstream tasks—including bidirectional forecasting, spatial imputation, and trajectory extrapolation—under a single generative score network.
- Minimalist Scaling and Inductive Regularization (UrbanFM): Chen et al. [61] present UrbanFM, scaling spatio-temporal foundation models up to 120M parameters. UrbanFM challenges over-engineered spatial graph heuristics, proving that minimalist self-attention with constrained inductive biases and MiniST tokenization (which normalizes variance across disparate city sensors) achieves superior scaling laws and cross-city zero-shot transfer compared to complex graph convolution-transformer hybrids.
- Satellite Image Time Series Foundation Models (TiMo): Qin et al. [62] introduce TiMo, scaling spatio-temporal foundation architectures to Satellite Image Time Series (SITS). Departing from standard vision transformers that process static frames independently, TiMo incorporates a

hierarchical spatio-temporal gyroscope attention mechanism that captures evolving multi-scale seasonal patterns across planetary terrain. Pre-trained on the MillionST dataset comprising one million satellite image phases across 100,000 geographic locations over five years, TiMo unifies deforestation monitoring, crop classification, and flood risk segmentation under a single masked autoencoding foundation.

- **Unified Conceptual Vision and Generalization Challenges:** Goodge et al. [57] formulate a comprehensive taxonomy for STFMs, identifying four orthogonal dimensions of generalization required for true foundational capability: (1) Spatial Generalization across varying graph topologies and geographic layouts; (2) Temporal Generalization across diverse sampling resolutions and non-stationary distribution shifts; (3) Scale Generalization across localized sensor clusters and regional macro-grids; and (4) Cross-Domain Generalization bridging human mobility, climate phenomena, and cyber-physical infrastructure.

These advances demonstrate that spatial graphs and 1D temporal series can be harmoniously unified within scalable transformer and diffusion backbones, establishing spatio-temporal modeling as a vibrant pillar of modern foundation model research.

4.8 Lightweight, Edge-Native, and Microcontroller TSFMs

While industrial-scale foundation models scale from hundreds of millions to billions of parameters, mission-critical operational technology (OT)—such as cyber-physical manufacturing loops, wireless telecommunications infrastructure, and embedded microcontrollers—operates under rigid constraints of sub-10ms response times, micro-watt energy budgets, and sub-100MB RAM ceilings. This has catalyzed a dedicated design movement toward edge-native foundation models:

- **Tiny-TSM (Birkel, 2025) [63]:** Birkel presents Tiny-TSM, establishing that foundation capabilities can be synthesized within an ultra-compact 23M parameter footprint. Trained on a single consumer-accessible A100 GPU in less than one week using the synthetic SynthTS generation pipeline, Tiny-TSM incorporates a novel causal input normalization scheme. This formulation enables dense next-token prediction loss without autoregressive rollout collapse, matching or outperforming industrial-scale billion-parameter foundation baselines across medium- and long-horizon forecasting tasks under mean squared error.
- **APEX: Network-Native Wireless Telemetry (Pradhan et al., 2026) [64]:** Developed at Cisco Systems, APEX addresses the failure mode where generic time series models falter on bursty, zero-inflated, and multi-layer coupled wireless network telemetry. APEX introduces a decoder-only transformer pre-trained on 10-channel telemetry spanning 4,500 production enterprise networks ($\approx 100K$ access point time series). It is released in two operational tiers: APEX-Large (269M parameters for cloud analytics, reducing MAE by 18% over Toto [45]) and APEX-Edge (10.5M parameters), which executes sub-second, privacy-preserving zero-shot anomaly detection directly on router/AP-class edge hardware.
- **Ultra-Lightweight Patch-Independent MLPs (Cheraghinia et al., 2025) [65]:** Cheraghinia et al. demonstrate that

transformer self-attention is unnecessary for constrained edge devices. By deploying simple patch-independent multi-layer perceptron (MLP) encoders comprising a mere 21K trainable parameters, their foundation architecture executes downstream wireless modulation classification, technology recognition, and line-of-sight detection in 0.33ms on commodity microcontrollers, maintaining $> 97\%$ accuracy across previously unseen data classes without representation degradation.

- **MACROCAST: Vintage-Consistent Real-Time Foundations (Carriero et al., 2026) [66]:** Carriero, Pettenuzzo, and Shekhar formulate MACROCAST, an ultra-efficient 15M-parameter foundation model for real-time macroeconomic forecasting that systematically eliminates both temporal lookahead contamination and revision bias. Standard TSFMs trained on fully revised time series suffer from hindsight leakage because historical macroeconomic indicators (e.g., GDP, non-farm payrolls) undergo extensive post-hoc revisions unavailable to real-time decision-makers. MACROCAST rules out both leakage forms by pretraining on purely synthetic time series generated from Bayesian Vector Autoregressions (BVARs) and Dynamic Factor Models (DFMs) in one GPU-day, followed by nine-minute fine-tuning on vintage-specific ALFRED releases. On the real-time FRED-MD benchmark, MACROCAST beats AR(1) on 80% of horizon pairs and surpasses Chronos-2 [8] and classical BVARs without data leakage.
- **TQS-PTQ: Dynamical Systems Quantization Sensitivity (Pavlova et al., 2026) [67]:** Pavlova et al. reframe post-training quantization (PTQ) for time-series foundation forecasters through the mathematical lens of dynamical systems stability. Conventional PTQ sensitivity heuristics are tightly coupled to specific quantizer bit-widths and calibration sets. In contrast, the Trajectory-based Quantization Sensitivity Score (TQS) models the autoregressive latent rollout as a discrete-time dynamical system $\mathbf{z}_{t+1} = \Phi(\mathbf{z}_t, \mathbf{x}_t)$, rigorously bounding how local quantization perturbation errors amplify across the forecasting horizon. Decoupled from quantizer selection, TQS-PTQ derives an exact sensitivity budget across layers prior to quantization, enabling robust mixed-precision allocation (FP16 down to INT4) for compiled, black-box edge models without requiring calibration data or second-order Hessian approximations.
- **Calibration Bets on the Past: 4-Bit Activation Calibration (Ye & Wanjiku, 2026) [68]:** Ye and Wanjiku conduct an extensive empirical audit of post-training quantization across 560 neural forecasters and eight walk-forward testing years on S&P 500 volatility series. Their findings demonstrate a fundamental dichotomy between weight and activation quantization: while 8-bit quantization is largely insensitive to calibration methodology, static 4-bit activation quantization causes severe degradation under standard absolute-maximum (abs-max) calibration, destroying 11% to 62% of predictive signal. Replacing abs-max with percentile calibration recovers 53% to 94% of this loss by effectively filtering temporal outliers and structural distribution breaks, establishing activation calibration as an indispensable prerequisite for 4-bit time-series deployment.
- **On-Device Fully Quantized Training on Microcontrollers (MCU-FQT, Deutel et al., 2024) [69]:** Deutel et al. break

through the barrier of edge adaptability on bare-metal ARM Cortex-M microcontrollers. Addressing the severe resource constraints of microcontrollers (< 256 KB SRAM and < 1 MB Flash memory), they formulate Fully Quantized Training (FQT), executing both forward inference passes and backward gradient propagation entirely within 8-bit integer arithmetic. Paired with dynamic partial gradient updates that selectively update only high-sensitivity layers, MCU-FQT enables embedded sensor forecasters to continuously fine-tune to streaming distribution shifts on-device without floating-point emulation overhead.

- Sparse Binary Transformers: 1-Bit Weights for Multivariate Time Series (Gorbett et al., 2023) [70]: Gorbett, Shirazi, and Ray pioneer extreme 1-bit binary parameterization ($\mathbf{W} \in \{-1, +1\}^{d_{\text{out}} \times d_{\text{in}}}$) for multivariate time-series transformers. Rather than relying on floating-point matrix multiplications $\mathbf{W}\mathbf{x}$, binary quantization transforms linear projections in multi-head self-attention and patch feedforward blocks into addition-only accumulation:

$$\mathbf{y} = \mathbf{W}\mathbf{x} = \sum_{j:W_{ij}=+1} x_j - \sum_{j:W_{ij}=-1} x_j. \quad (15)$$

Coupled with learnable binary scaling vectors and multi-rate patch downsampling, Sparse Binary Transformers achieve an 87% parameter compression ratio and an order-of-magnitude reduction in arithmetic logic unit (ALU) gate toggle rates on edge FPGAs and embedded ASICs, proving that temporal attention maps can be effectively learned under binary weight discretization.

- Q-DEQ: Quantized Deep Equilibrium Forecasters (Yang et al., 2026) [71]: Yang et al. address edge deployment coding constraints by bridging Deep Equilibrium Models (DEQs) with low-bit quantization. Conventional deep architectures stack explicit feedforward layers, incurring memory footprints proportional to network depth. In contrast, DEQs simulate infinite effective depth by computing the implicit fixed-point equilibrium state:

$$\mathbf{z}^* = f_{\theta}(\mathbf{z}^*, \mathbf{x}), \quad (16)$$

using a single shared transformation layer (0.3M parameters). Q-DEQ resolves the numerical divergence that standard fixed-point Broyden solvers experience under discrete INT4/INT8 arithmetic by formulating a discrete-solving relaxation and contraction-mapping boundary constraint. Q-DEQ achieves parity with full-precision FP32 PatchTST and DLinear baselines while cutting memory storage to 75 KB, enabling continuous dynamical equilibrium modeling within microcontroller cache bounds.

- SGA: State-Guided Autoregressive Uncertainty Calibration (Hu et al., 2026) [72]: Hu et al. address the uncalibrated error accumulation inherent to long-horizon autoregressive rollouts in foundation forecasters. While multi-step forecasting models recursively feed predicted patches $\hat{\mathbf{y}}_t$ back as context $\mathbf{x}_{t+1}$, epistemic errors cascade exponentially, producing miscalibrated predictive intervals. SGA formulates a State-Guided Autoregressive adapter: by extracting the foundation backbone’s internal hidden transition trajectories $\mathbf{h}_t$ and conditioning a dual-headed Gaussian-process parametric head on state-space curvature, SGA dynamically scales rollout variance without requiring com-

putationally prohibitive Monte Carlo sampling or ensemble inference.

These architectures prove that foundation pretraining is not intrinsically tied to massive parameter scales; structured inductive priors, synthetic data curation, and patch-level modularity enable sub-megabyte foundation models for extreme edge environments.

4.9 Non-Stationary Temporal Drift and Test-Time Adaptation (TTA)

A foundational vulnerability of all pre-trained temporal forecasters is the assumption of distributional stationarity between pretraining corpora $\mathcal{D}_{\text{pre}}$ and streaming inference horizons $\mathcal{D}_{\text{test}}$. In open cyber-physical, meteorological, and economic environments, time series are intrinsically non-stationary, exhibiting evolving mean-variance trajectories, seasonal structural breaks, and abrupt concept drift:

$$\mathcal{P}_{\text{train}}(\mathbf{X}_{t-L:t}) \neq \mathcal{P}_{\text{test}}(\mathbf{X}_{t-L:t}), \quad \mathcal{P}_{\text{train}}(\mathbf{Y}_{t:t+H} \mid \mathbf{X}_{t-L:t}) \neq \mathcal{P}_{\text{test}}(\mathbf{Y}_{t:t+H} \mid \mathbf{X}_{t-L:t}) \quad (17)$$

While standard fine-tuning requires labeled supervision and risks catastrophic forgetting of general priors, Test-Time Adaptation (TTA) enables pre-trained foundation models to continually adapt their latent representations to streaming distribution shifts using only unlabeled observations available at inference time [73], [74].

Recent advances formalize TTA for foundation time-series forecasters across three primary paradigms:

- Online Test-Time Adaptation Forecasters (TSF-TTA): Kim et al. [73] establish the TSF-TTA framework, systematically addressing non-stationarity in temporal models without requiring access to source pretraining datasets. TSF-TTA freezes the feature extraction backbone and restricts test-time updates to reversible instance normalization statistics and lightweight dynamic linear adapter heads. By minimizing an unsupervised temporal consistency objective across sliding test windows, TSF-TTA dynamically tracks concept drift, cutting test-time prediction error by up to 22% across non-stationary benchmarks while incurring negligible compute overhead.
- Continuous Dynamical Adaptation via Neural ODEs (AdaNODEs): Dang et al. [74] introduce AdaNODEs, moving beyond discrete heuristic gradient steps by formulating test-time parameter trajectories through continuous-time Neural Ordinary Differential Equations (NODEs). In AdaNODEs, the evolution of latent state representations $\mathbf{z}(t)$ under streaming temporal drift is governed by a parameterized continuous vector field:

$$\frac{d\mathbf{z}(t)}{dt} = f_{\theta_{\text{ode}}}(\mathbf{z}(t), t), \quad \mathbf{z}(t_1) = \mathbf{z}(t_0) + \int_{t_0}^{t_1} f_{\theta_{\text{ode}}}(\mathbf{z}(\tau), \tau) d\tau. \quad (18)$$

By updating only an ultra-compact 0.15M parameter NODE adapter rather than backbone weights, AdaNODEs smoothly accommodates irregular sampling intervals and continuous distribution shifts, demonstrating superior parameter efficiency and convergence stability over standard gradient descent.

- Regime-Guided Meta-Control (RG-TTA): Kumar et al. [75] resolve the dilemma between under-adaptation (failing to catch sudden shifts) and over-adaptation (overfitting to

transient noise). RG-TTA introduces a regime-guided meta-controller that continuously computes rolling distribution statistics (spectral energy, mean drift, variance change) to classify the streaming series into stationary baseline vs. shifted regimes. The meta-controller dynamically modulates the adaptation learning rate $\eta_t = \sigma(\mathbf{W}_r \mathbf{r}_t + b_r)$ and gradient clipping thresholds, applying aggressive updates during genuine regime changes while damping updates during stable periods to preserve foundational knowledge.

4.10 Causal Discovery, Structural Priors, and Counterfactual Foundation Models

While contemporary time series foundation models exhibit remarkable correlational forecasting capabilities, their black-box representations struggle with out-of-distribution intervention reasoning, policy simulation, and root-cause attribution [76]. Integrating Structural Causal Models (SCMs) into foundation architectures represents a crucial frontier bridging deep representations with causal mechanics [77], [78].

Formally, a dynamic temporal structural causal model defines a multivariate time series $\mathbf{X}_t = [X_{1,t}, \dots, X_{D,t}]^\top$ across time steps t via a system of structural equations over parent variables $\text{PA}(X_{i,t})$ and independent exogenous noise processes $U_{i,t}$:

$$X_{i,t} := f_i(\text{PA}_{\text{inter}}(X_{i,t}), \text{PA}_{\text{intra}}(X_{i,t}), U_{i,t}), \quad U_{i,t} \sim \mathcal{P}_{U_i}, \quad (19)$$

where $\text{PA}_{\text{inter}}(X_{i,t}) \subseteq \{X_{j,t-\tau} \mid j \in [D], \tau \geq 1\}$ denotes lagged inter-temporal causal parents and $\text{PA}_{\text{intra}}(X_{i,t}) \subseteq \{X_{j,t} \mid j \neq i\}$ represents instantaneous intra-temporal parents, defining an underlying causal Directed Acyclic Graph (DAG) $\mathcal{G}_{\text{causal}}$.

Recent foundation architectures synthesize causality across three complementary frontiers:

- **Foundation Models for Time-Series Causal Discovery (Causal-PT):** Stein et al. [77] demonstrate that causal DAG discovery from temporal signals can be learned end-to-end by a pre-trained foundation transformer. Rather than relying on combinatorial constraint-based testing (e.g., PCMLCI) or continuous optimization (e.g., DYNOTEARS), Causal-PT pre-trains a 25M-parameter transformer over 10 million synthetically generated nonlinear structural causal processes. At test time, Causal-PT maps raw multivariate temporal trajectories directly to the estimated causal adjacency matrix $\hat{\mathbf{A}}$ in a single feedforward pass, outperforming classical algorithms while exhibiting high robustness to observational noise.
- **Interventional Time Series Priors for In-Context Causal Models (CausalTimePrior):** Thumm and Chen [78] pioneer the fusion of Prior-Data Fitted Networks (PFNs) with temporal structural causality. Prior PFNs (e.g., TabPFN, ForecastPFN) were constrained strictly to observational data. CausalTimePrior constructs a synthetic generator producing hundreds of millions of diverse temporal structural causal models (TSCMs) paired with both observational trajectories and atomic interventions $do(X_j = \alpha)$. By pre-training a transformer on these paired trajectories, CausalTimePrior performs exact in-context causal effect estimation and counterfactual forecasting in a single forward evaluation without requiring gradient optimization.

- **Causal Time Series Generation via Diffusion Models (CaTSG):** Xia et al. [79] introduce CaTSG, establishing a structural causal diffusion framework for multivariate time series. CaTSG incorporates time-varying causal graph representations into the reverse-time score-matching process:

$$s_\theta(\mathbf{x}_t, t, \mathcal{G}_{\text{causal}}) \approx \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t \mid \mathcal{G}_{\text{causal}}). \quad (20)$$

By constraining the generative diffusion trajectories through learned causal adjacency masks, CaTSG guarantees that generated trajectories strictly adhere to physical causal dependencies, enabling controllable simulation across observational, interventional, and counterfactual query spaces.

4.11 Continuous-Time Asynchronous Event Streams and Market Microstructure Foundations

A fundamental limitation of conventional time series foundation models lies in their rigid presumption of uniform, synchronous discrete sampling grids ($\Delta t = \text{const}$). In real-world environments—such as financial limit order books (LOBs), electronic health records, and seismic sensor networks—observations arrive as continuous-time asynchronous event streams $\mathcal{S} = \{(t_i, \mathbf{m}_i)\}_{i=1}^N$, where inter-arrival intervals $\tau_i = t_i - t_{i-1} \in (0, \infty)$ span multiple orders of magnitude and marks $\mathbf{m}_i \in \mathcal{M}$ encode heterogeneous event features. Imposing artificial uniform binning incurs substantial temporal distortion, aliasing high-frequency bursts while creating sparse, computationally wasteful zero-padded windows during quiescent periods.

Formally, an asynchronous event sequence is governed by a marked temporal point process characterized by a conditional intensity function $\lambda^*(t) = \lim_{\Delta t \rightarrow 0} \frac{\mathbb{P}(\text{event in } [t, t+\Delta t] \mid \mathcal{H}_t)}{\Delta t}$, where $\mathcal{H}_t$ represents the filtration of historical events up to time t . Foundation modeling over such continuous-time streams requires architectural invariants capable of handling arbitrary time deltas and cross-dataset intensity scale mismatches:

- **SurF: Generative Continuous-Time Foundation Models via Time Rescaling** (Rezaei et al., 2026) [80]: Rezaei, Balaji, and Krishnan resolve the intractable quadrature integration problem in neural temporal point processes by leveraging the classical Time Rescaling Theorem (TRT) as a learnable bijective continuous flow. The TRT establishes that any point process with continuous cumulative intensity $\Lambda(t) = \int_0^t \lambda^*(s) ds$ can be transformed into a standard unit-rate Poisson process:

$$\xi_i = \Lambda(t_i) - \Lambda(t_{i-1}) \stackrel{\text{i.i.d.}}{\sim} \text{Exp}(1). \quad (21)$$

SurF parameterizes $\Lambda_\theta(t)$ via monotonic neural networks and inverts this transformation: during training, event inter-arrival times are mapped into standardized exponential noise, while generative sampling draws $\xi_i \sim \text{Exp}(1)$ and computes $t_i = \Lambda_\theta^{-1}(\Lambda_\theta(t_{i-1}) + \xi_i)$ in closed form. Crucially, mapping heterogeneous streams to a unified unit-rate latent space enables multi-dataset joint pretraining across earthquakes, retweets, and e-commerce streams, eliminating loss scale incompatibilities and establishing the first scalable foundation model for continuous irregular event forecasting.

- **TradeFM: Generative Market Microstructure Foundation Models** (Kawawa-Beaudan et al., 2026) [81]: Developed at

J.P. Morgan AI Research, TradeFM introduces a 524M-parameter generative transformer foundation model trained on billions of asynchronous order-flow events spanning over 9,000 global equities. TradeFM addresses the severe cross-asset non-stationarity of financial microstructure by introducing scale-invariant universal tokenization: order book depths, trade volumes, and micro-price updates are mapped into normalized continuous event embeddings decoupled from idiosyncratic stock price scales. Coupled with a high-fidelity deterministic stock market matching engine, TradeFM performs zero-shot generative simulation of market dynamics, accurately replicating non-linear stylized facts (volatility clustering, heavy-tailed return distributions, and order-flow autocorrelation decay) while generalizing zero-shot across distinct international geographic regimes (e.g., European and APAC exchanges).

4.12 Physics-Informed Foundations and Scientific Constraint Injection

In mission-critical physical and industrial domains—such as electric power grids, aerospace aerodynamics, and renewable energy generation—purely data-driven foundation models face severe adoption barriers: their unconstrained predictions frequently violate invariant physical conservation laws (e.g., Kirchhoff’s laws, energy conservation, power balance constraints), leading to catastrophic failures if deployed in autonomous control loops. Embedding physical inductors directly into foundation architectures and pretraining pipelines represents an essential paradigm for verifiable scientific foundation models:

- GridSFM: Physics-Inspired Foundation Models for Power Flow (Bhan et al., 2026) [82]: Bhan et al. develop GridSFM, a 15M-parameter physics-inspired foundation model for solving AC Optimal Power Flow (AC-OPF) across large-scale electrical transmission grids. Rather than treating power flow as generic spatio-temporal regression, GridSFM embeds the non-linear AC power balance equations:

$$S_i = V_i \sum_{k \in \mathcal{N}_i} Y_{ik}^* V_k^*, \quad \forall i \in \mathcal{V}_{\text{bus}}, \quad (22)$$

directly into a physics-guided graph neural network. To address non-convex feasibility boundaries and disconnected solution manifolds, GridSFM introduces a lifted problem reformulation with logarithmically penalized slack variables. Pretrained over 54 distinct transmission topologies (500 to 4,000 buses), GridSFM achieves a remarkable 2.45% zero-shot generation-cost error when evaluated on unseen 10,000-bus power grids, providing warm-start trajectories that accelerate commercial interior-point solvers by over an order of magnitude.

- Physics-Informed Synthetic Histories for Cold-Start Forecasting (Longarini et al., 2026) [83]: Longarini et al. tackle the severe “cold-start” barrier in renewable energy forecasting, where newly commissioned solar photovoltaic (PV) facilities lack historical telemetry required for foundation model in-context conditioning. The authors introduce a two-stage physics-informed pipeline: deterministic physical model chains (incorporating solar positioning, irradiance transposition, and thermal cell equations via pvlib) construct synthetic production histories $\tilde{\mathbf{X}}_{1:T}$ solely from geographical coordinates and numerical weather predictions.

Feeding these physically consistent synthetic histories into off-the-shelf foundation models (TabPFN-TS, Chronos-2, TimesFM) without fine-tuning reduces forecasting error (MAE) by up to 50% compared to classical statistical models across 440 commercial installations, demonstrating that physics-based surrogate conditioning bridges data scarcity in zero-shot foundation deployments.

5 Repurposed Large Language Models (LLM4TS)

In parallel to native foundation models, an active branch of literature investigates repurposing existing Large Language Models (LLMs) for time series analysis. Driven by the hypothesis that multi-billion parameter text models have acquired generalized pattern matching, sequence extrapolation, and relational reasoning abilities, these methods adapt frozen or fine-tuned language backbones to temporal streams.

5.1 Cross-Modal Token Reprogramming

Token reprogramming bridges the structural gap between continuous numerical signals and the discrete token space of language transformers by learning parameter-efficient adapter layers while keeping the core language model frozen.

5.1.1 Time-LLM: Prompt-Guided Reprogramming

Introduced by Ming Jin et al., Time-LLM [19] establishes a canonical reprogramming framework. Input time series are patched into subseries vectors, which are mapped through a linear projection layer into the textual word embedding space of Llama-7B. Crucially, Time-LLM constructs dynamic text prompts containing domain context, task instructions, and dataset statistics (e.g., trend, variance, seasonality). The frozen language model processes the concatenated prompt tokens and reprogrammed temporal tokens, producing representations that are decoded back into continuous forecasts via a linear projection head. Time-LLM demonstrates strong zero-shot and few-shot capabilities across benchmark datasets.

5.1.2 One Fits All (GPT4TS / OFA): Universal Backbone Sharing

In One Fits All (GPT4TS) [20], Zhou et al. explore the minimal adaptation required to transfer language transformers to temporal tasks. By taking a pre-trained GPT-2 backbone, freezing its self-attention and feed-forward layers, and retraining only an input linear patch embedder and output layer, GPT4TS achieves state-of-the-art results across forecasting, anomaly detection, classification, and imputation. The authors argue that multi-head self-attention trained on natural language acts as a universal sequence engine capable of recognizing temporal autocorrelation and periodic waveforms.

5.1.3 Decomposition, Prototype Alignment, and Distillation

Subsequent works systematically refine cross-modal reprogramming:

- TEST [31]: Sun et al. identify the semantic gap between continuous temporal tokens and discrete word embeddings. TEST constructs a learnable dictionary of text prototypes derived from language vocabulary clusters. By mapping

temporal patches onto convex combinations of text prototypes, TEST aligns time-series patterns with linguistic semantics before feeding them into the frozen LLM.

- CALF [32]: Liu et al. address catastrophic representation degradation during cross-modal adaptation. CALF introduces Cross-modal Alignment and Low-rank Fine-tuning, utilizing word-level and sentence-level semantic distillation to preserve the backbone’s latent sequence modeling capabilities while aligning temporal features.
- UniTime [33]: Liu et al. design a language-empowered unified model for cross-domain forecasting. By introducing domain-specific language instructions and dynamic masking, UniTime transfers a single LLM across multiple distinct data distributions without catastrophic interference.
- LLM4TS [84]: Chang et al. propose a two-stage alignment strategy: a self-supervised time-series alignment pre-training phase followed by downstream forecasting fine-tuning, coupled with a two-level temporal aggregation module to capture multi-scale seasonality.
- TimeCMA [85]: Chen et al. design cross-modality alignment between multivariate channels and language representations, using channel-wise text prompts to guide inter-series dependency modeling.
- TEMPO [21]: Decomposes input series into trend, seasonal, and residual components using classical STL decomposition. Each component is independently embedded and guided by specialized prompts, enabling the language model to reason over modular temporal dynamics.
- S²IP-LLM [86]: Introduces semantic-space informed prompting, projecting statistical temporal attributes into dense semantic embeddings that constrain the LLM’s generative search space.
- LLM-Mixer [87]: Kowsher et al. formulate multiscale token-mixing inside frozen language models. By decomposing temporal signals across multiple sampling strides and processing them via parallel attention pathways in LLaMA-2, LLM-Mixer captures multiscale periodicities without full-parameter adaptation.
- TAC-Time [88]: Liang et al. propose the “Texts as Channels” architecture for multimodal forecasting. Instead of forcing numerical values into linguistic strings or shared embedding spaces, TAC-Time treats contextual text streams as auxiliary multivariate channels, employing asymmetric dual-stream cross-attention to inject semantic knowledge into temporal representations without disrupting continuous numerical geometry.

5.2 Direct Prompting and Decimal Tokenization

Rather than learning adapter layers, direct prompting methods interact with off-the-shelf closed or open LLMs (such as GPT-4 or Llama-3) by formatting numeric series as text strings.

5.2.1 Foundations: PromptCast and Generative Formulation

The paradigm of casting numerical forecasting as text generation was pioneered by Xue and Salim in PromptCast [42]. Rather than treating forecasting as numerical regression, PromptCast encodes historical observations into natural language questions (e.g., “Given the temperature readings of the

past 24 hours, what will be the temperature over the next 6 hours?”) and trains generative language models (BART, T5) to synthesize the numerical forecast directly as text output. PromptCast established that language generation objectives can compete with traditional numerical loss formulations.

5.2.2 Decimal String Modeling: LLMTIME

In LLMTIME [6], Gruver et al. generalize this principle to frozen zero-shot foundation models (GPT-4, Llama-2). By formatting time series as comma-separated ASCII digit strings (e.g., “12.4, 15.8, 14.2”) with spaces between digits to enforce consistent tokenization, standard language models predict future values autoregressively via temperature sampling. This enables probabilistic forecasting without any fine-tuning.

However, direct decimal prompting incurs severe practical drawbacks:

- 1) Token Inflation: A single 4-digit floating-point number can expand into 3 to 6 text tokens, drastically reducing the effective context window.
- 2) Inference Latency: Autoregressively sampling hundreds of tokens from a 70B parameter model for a single univariate series takes several seconds, rendering real-time streaming infeasible.
- 3) Financial Cost: Using proprietary commercial API endpoints (e.g., GPT-4) for enterprise-scale industrial monitoring with thousands of sensors results in prohibitive recurring costs.

5.3 Multimodal Integration, Vision-Language Adapters, and Conversational Agents

Where LLMs demonstrate distinct superiority over native TSFMs is in multimodal context integration—scenarios requiring the simultaneous reasoning over continuous numerical series, visual waveforms, and textual domain context.

5.3.1 Visual Cross-Modal Forecasters

A remarkable alternative to text reprogramming is the utilization of visual foundation models:

- VisionTS [34]: Chen et al. uncover that visual masked autoencoders (e.g., MAE pretrained on ImageNet) act as “free-lunch” zero-shot forecasters. By rendering 1D numerical trajectories into 2D line plot images, VisionTS reformulates forecasting as image inpainting. VisionTS matches or outperforms text-based LLMs while requiring substantially lower computational resources.
- Time-VLM [35]: Zhong et al. explore multimodal Vision-Language Models (e.g., LLaVA, CLIP) for augmented forecasting. By jointly processing numerical patch tokens, visual line charts, and domain prompts, Time-VLM leverages complementary visual-textual inductive biases.

5.3.2 Conversational Agents and Omni-Modal Intelligence

Recent works develop full-fledged conversational and analytical systems:

- ChatTS [89]: Aligns time series with LLMs using synthetic conversational datasets generated through rule-based simulations, enabling grounded interactive Q&A, anomaly explanation, and trend narration.
- ChatTime [90]: A unified multimodal foundation model that bridges numerical waveforms and textual queries for

interactive dialogue, counterfactual reasoning, and zero-shot forecasting.

- Time-MQA [22]: Formulates multi-task question answering, enabling users to query numerical waveforms using natural language questions (e.g., “Did the turbine vibration exceed safety thresholds after 14:00?”).
- TimeOmni-1 & TimeOmni-VL [23], [91]: Advance multimodal temporal foundation modeling. TimeOmni-1 generates numeric forecasts alongside chain-of-thought rationales, while TimeOmni-VL unifies continuous time series, text, and visual spectrograms into an omni-modal transformer.
- TRACE [92]: Kan et al. propose temporal conditional estimation for multimodal foundation models, formalizing joint inference over complex clinical physiological waveforms, clinical electronic health record (EHR) text narratives, and diagnostic interventions under non-stationary physical dynamics.
- OpenTSLM [93]: Specializes multimodal reasoning for clinical healthcare, integrating multi-channel ICU physiological waveforms (ECG, arterial blood pressure) with clinical notes and physician orders.

5.3.3 Spatio-Temporal Large Language Models

Extending LLM cross-modal adaptation from 1D sequences to interconnected geographic networks represents an impactful frontier:

- UrbanGPT [94]: Li et al. (KDD 2024) introduce a spatio-temporal large language model for urban forecasting. UrbanGPT constructs a Spatio-Temporal Dependency Encoder (ST-Encoder) that maps spatial topological graphs and historical temporal dynamics into language token representations. When interfaced with a frozen LLaMA-2-7B backbone via lightweight cross-modal instruction tuning, UrbanGPT generalizes zero-shot across diverse urban prediction tasks (traffic flow, ride-hailing demand) in unseen metropolitan regions where training data is scarce.
- ST-LLM [95]: Liu et al. (TKDE 2025) develop a spatial-temporal large language model framework for traffic prediction. ST-LLM integrates spatial graph positional embeddings with partially-frozen transformer attention layers, enabling language models to simultaneously capture localized spatial dependencies across physical sensor networks and long-range temporal trends without incurring full model fine-tuning overhead.

5.4 Omni-Modal Grounding and Vision-Language-Time Pretraining

While initial multimodal efforts relied on pipelined adapters or frozen post-hoc alignments, the cutting edge of cross-modal temporal research focuses on omni-modal grounding—binding continuous physical sensor waveforms, 2D visual spectrograms/imagery, and discrete textual semantic narratives into unified perceptual representations:

- VisionTS++: Continual Pre-trained Vision Backbones [96]: Shen et al. significantly advance visual time series modeling over vanilla VisionTS. Recognizing that off-the-shelf vision transformers pre-trained solely on natural images suffer from scale mismatch and domain discrepancy when processing line plots, VisionTS++ introduces multi-scale

image projections and a continual pre-training curriculum on diverse synthetic and cross-domain time-series line reconstructions. This continual adaptation enables the vision backbone to capture multiscale cyclical waveforms and subtle trend deviations with superior fidelity.

- VLT: Industrial Multimodal Intelligence [97]: Wang et al. formulate VLT, a vision-language-time multimodal foundation architecture for complex industrial Prognostics and Health Management (PHM). Recognizing the semantic chasm between raw numerical sensor vibrations and diagnostic text manuals, VLT deploys the time-frequency spectrogram as a bidirectional visual bridge. It incorporates a Time-aware Mixture-of-Experts (Time-MoE) to handle heterogeneous physical dynamics alongside a Frequency-Text Augmented Learner operating in a shared semantic manifold, governed by a time-centric gradient alignment mechanism that eliminates cross-modal optimization conflicts during multi-task pretraining.
- Chronicle: Joint Pretraining From Scratch [98]: Quinlan et al. challenge the prevailing dogma of post-hoc LLM adaptation by introducing Chronicle, a compact 324M parameter decoder-only transformer trained ab initio jointly on natural language corpora and time series within a single homogeneous architecture. By sharing transformer layers, self-attention heads, and residual streams across interleaved text tokens and subseries patches, Chronicle proves that genuine cross-modal representation emerges directly from shared pretraining. Chronicle matches Gemma-3-270M-PT on 19 natural language understanding (NLU) benchmarks while simultaneously outperforming supervised fusion baselines on the multimodal Time-MMD benchmark and setting new state-of-the-art accuracy on frozen-embedding UCR/UEA classifications.
- ChronoSteer: Synthetic Cross-Modal Steering [99]: Addressing the acute scarcity of real-world paired text-time corpora, Wang et al. propose ChronoSteer. ChronoSteer deploys an agentic framework where an LLM translates narrative contextual events into structured revision instructions that dynamically modulate the zero-shot forecasts of a frozen numerical TSFM (e.g., Chronos-Bolt). By discretizing continuous revisions into a compact codebook of instruction anchors, ChronoSteer achieves a 25.8% accuracy enhancement over standalone unimodal backbones without requiring end-to-end retraining.

These developments indicate that while native TSFMs dominate pure high-throughput numerical prediction, omni-modal foundation architectures and repurposed multi-modal LLMs serve as indispensable cognitive orchestrators for real-world cyber-physical intelligence.

5.5 Interactive Time-Series Agents, Tool-Augmented Reasoning, and Autonomous Exploration

The conventional paradigm in deep temporal modeling treats time series forecasting as a static, single-pass mapping:

$$\hat{\mathbf{Y}}_{t:t+H} = f_{\theta}(\mathbf{X}_{t-L:t}), \quad (23)$$

where a frozen neural network ingests fixed-length historical observations and produces future trajectories in isolation. However, in real-world cyber-physical systems, financial trading desks, and industrial operations, time series analysis

is intrinsically dynamic, exploratory, and multi-turn [100]. Decision-makers require interactive systems capable of testing scientific hypotheses, cleaning corrupted signals, invoking specialized analytical tools, and explaining anomalous behaviors through multi-turn dialogue.

To bridge this operational chasm, recent research pivots toward interactive time-series agents that integrate LLM-based reasoning backbones with tool execution engines and external memory banks:

- **Real-Time Streaming Time-Series Interaction (TimeInteract):** Pan et al. [101] uncover a critical limitation of existing time-series language models: their static nature requires either offline batch sequences or halts incoming sensor ingestion while generating autoregressive text. Pan et al. introduce TimeInteract, pioneering a real-time streaming interaction regime where the agent continuously perceives asynchronous temporal streams, dynamically decides when to remain silent or respond, and continues processing incoming telemetry during token generation. TimeInteract deploys a dual-view streaming encoder (capturing local transient variations and long-term historical trends) alongside a decoupled streaming inference engine. Evaluated on StreamTSI-34K (comprising 34,588 streaming episodes and 77,505 interactive responses), TimeInteract achieves near-zero stream stall and a $2.15\times$ inference speedup while improving response trigger accuracy across hierarchical interaction levels.
- **Tool-Augmented Sequential Decision Policies (Cast-R1):** Tao et al. [102] reformulate time series forecasting from a static mapping into a sequential decision-making problem. Cast-R1 introduces a memory-based state management mechanism that maintains intermediate analytical evidence across interaction turns. The agent autonomously navigates a modular toolkit: extracting seasonal-trend decompositions, running statistical stationarity checks, invoking lightweight forecasting models (e.g., DLinear, PatchTST) for candidate predictions, and executing iterative self-reflection to revise forecasts. Cast-R1 is trained through a two-stage curriculum combining supervised fine-tuning (SFT) with multi-turn reinforcement learning (RL), setting superior accuracy on complex non-stationary series.
- **Agentic Time Series Reasoning and TimeToolBench (TimeART):** Wu et al. [103] address Time Series Question Answering (TSQA) by creating TimeART. Recognizing that generic LLMs fail to manipulate continuous numerical series with computational precision, TimeART fuses external data science tools with an 8B Time Series Reasoning Model (TSRM). The model is trained on TimeToolBench—a curated dataset of 100,000 expert analytical trajectories—via a four-stage curriculum that incorporates experience replay and self-reflection, establishing state-of-the-art multi-step numerical and causal reasoning.
- **Domain-Oriented Inference Agents (TS-Reasoner):** Ye et al. [104] present TS-Reasoner, a domain-specialized agent combining LLM symbolic reasoning with domain-specific computational routines and an iterative error feedback loop. When evaluated on diagnostic benchmarks such as TimeSeriesExam, TS-Reasoner successfully overcomes the hallucinations and arithmetic inaccuracies of general-purpose LLMs, automating multi-step analytical workflows.

- **Data-Centric AutoML Agents (DCATS):** Yeh et al. [105] explore data-centric time series engineering. Rather than searching for increasingly complex model architectures, DCATS deploys an LLM planner to reason over time-series metadata, automatically diagnosing data corruption, imputing sensor dropouts, and optimizing preprocessing pipelines. DCATS achieves a 6% error reduction across diverse forecasting backends on massive traffic datasets.
- **Scientific Inquiry and Multi-Model Arbitration (TimeAgent):** Wang et al. (IEEE TKDE 2026) introduce TimeAgent, a training-free framework that addresses the empirical “seesaw effect” where individual TSFMs excel in specific domains while degrading on others. TimeAgent structures time series analysis as a closed-loop scientific inquiry: autonomously profiling data characteristics, formulating experiment plans, arbitrating predictions across heterogeneous foundation models (Chronos, TimesFM, Moirai), and translating quantitative outputs into interpretable natural language rationales.

These agentic systems indicate that the future of temporal AI lies not in monolithic forecasters, but in modular ecosystems where foundation models serve as specialized tools orchestrated by reasoning agents.

5.6 Multi-Agent Collaborative Ensembles, Swarm Deliberation, and Dynamic Consensus

While single-agent frameworks augment LLMs with external tools, they remain bounded by the cognitive blindspots of a solitary reasoning backbone. In real-world cyber-physical environments, financial trading desks, and macroeconomic planning, decision-making relies on interdisciplinary committees where domain specialists debate alternative hypotheses, challenge assumptions, and reconcile contradictory evidence [106]. To emulate this collective intelligence, recent research formalizes multi-agent collaborative ensembles and swarm deliberation for temporal reasoning and forecasting:

Let $\mathcal{A} = \{A_1, A_2, \dots, A_M\}$ denote a heterogeneous swarm of M specialized agents (e.g., visual trend analyzers, numerical statistical auditors, causal graph reasoners, and domain context interpreters). Rather than generating a solitary prediction in a single forward pass, the swarm engages in an iterative multi-turn deliberation:

$$m_i^{(k)} = \mathcal{F}_i \left(m_i^{(k-1)}, \mathcal{M}_{-i}^{(k-1)}, \mathbf{X}_{t-L:t} \right), \quad \mathcal{M}_{-i}^{(k-1)} = \{m_j^{(k-1)}\}_{j \neq i}, \quad (24)$$

where $m_i^{(k)}$ represents agent A_i ’s critique and refined rationale at deliberation round k . Deliberation proceeds until a termination criterion is met (e.g., semantic consensus or variance stabilization), whereupon a consensus aggregator synthesizes the finalized forecast $\hat{\mathbf{Y}}_{t:t+H}$:

$$\hat{\mathbf{Y}}_{t:t+H} = \text{Arbiter} \left(\{m_i^{(K)}\}_{i=1}^M, \mathbf{X}_{t-L:t} \right). \quad (25)$$

Key architectures advancing multi-agent temporal collaboration include:

- **Multimodal Collaborative Debate (MC-Debate, Trirat et al., 2026) [107]:** Trirat et al. identify an empirical paradox in cross-modal temporal reasoning: while visual representations convey global waveform geometry, they lack numerical precision; conversely, raw numerical patches provide point-level fidelity but obscure high-level semantic trends. To

resolve this tension, MC-Debate introduces a multi-agent debate framework comprising a Visual Specialist (ingesting line plot renderings), a Numerical Specialist (evaluating patch statistics and Fourier transforms), and a Textual Specialist (interpreting domain metadata). Coordinated by an impartial Arbiter Agent over three debate rounds, MC-Debate achieves a 14.2% error reduction across zero-shot reasoning benchmarks and eliminates 76% of hallucinated trend reversals compared to single-agent baselines.

- **Failure-Driven Self-Evolution of Time Series Agents (TimeEvo, Yang et al., 2026) [108]:** Yang et al. address the fundamental fragility of static tool sets. Traditional agents are endowed with fixed API libraries at initialization; when confronting unprecedented signal distortions or novel mathematical transforms, they experience persistent execution failures. TimeEvo introduces failure-driven self-evolution. Whenever an analytical trajectory encounters a task failure, a Meta-Evolutionary Agent inspects the error trace, autonomously codes a new Python tool function, verifies its correctness on synthetic counterfactual test cases, and dynamically registers the verified routine into the shared agent skill repository. Evaluated on TimeQA, TimeEvo’s evolving skill set continuously elevates task success rates without manual human intervention.
- **Self-Evolving Policies for Agentic Time-Series Forecasting (Hu et al., 2026) [109]:** Hu et al. formalize agentic forecasting under non-stationary environmental dynamics, where the relative efficacy of statistical algorithms, deep foundation models, and prompt strategies shifts across temporal regimes. Rather than relying on static prompt templates, Hu et al. model the agent’s arbitration strategy as a parameterized decision policy. As future observations materialize, the agent computes the realization regret and executes policy gradient updates, dynamically shifting its delegation weights between foundation backbones (e.g., Chronos vs. TimesFM) and localized autoregressive models in response to changing market volatility regimes.
- **Traceable Multi-Agent Systems for Knowledge-Based Forecasting (Kang et al., 2026) [106]:** Kang et al. tackle the black-box opacity that impedes enterprise adoption of foundation forecasters. Traceable Multi-Agent decomposes complex enterprise forecasting workflows into four verifiable sub-agents: Document Harvester (extracting supply-chain disruptions from regulatory 10-K filings), Data Transformer, Code Execution Engine, and Decision Traceability Auditor. The system maintains an immutable Directed Acyclic Graph (DAG) linking every numerical projection back to verified textual citations and executed code blocks, providing legal auditability and regulatory compliance for high-stakes financial operations.
- **MetaCaster: Meta-Harness Few-Shot Forecaster Agents (Shen et al., 2026) [110]:** Recognizing that deploying multi-billion-parameter foundation models is economically prohibitive in edge and low-power IoT scenarios, Shen et al. propose MetaCaster. MetaCaster employs an LLM-driven meta-harness agent that analyzes statistical fingerprint vectors (seasonality strength, spectral entropy, trend linearity) of target datasets to autonomously architect and train compact, specialized forecasters (<1M parameters) within minutes. By optimizing few-shot adaptation recipes and data augmentation schedules, MetaCaster matches the

zero-shot accuracy of 100M+ parameter foundation models while slashing inference compute by up to 95%.

These advances demonstrate that multi-agent deliberation, self-evolutionary tool creation, and traceable consensus represent a vital paradigm shift—transforming time-series intelligence from solitary pattern matchers into collaborative, self-improving scientific societies.

5.7 Reinforcement Learning from Human/Physical Feedback and Post-Training Policy Optimization

While Supervised Fine-Tuning (SFT) has served as the default adaptation mechanism for foundation models, SFT suffers from severe intrinsic pathologies in temporal domains: (1) exposure bias, where compounding autoregressive rollout errors diverge rapidly from the training distribution; (2) indifference to asymmetric business or physical costs, where an under-forecast penalty (e.g., hospital bed shortage, supply chain stockout) may be order-of-magnitude worse than an equivalent over-forecast; and (3) inability to elicit verifiable multi-step reasoning. Post-training policy optimization via Reinforcement Learning from Human Feedback (RLHF) and Reinforcement Learning from Physical Feedback (RLPF) bridges this critical divide.

Formally, consider a policy $\pi_\theta(\mathbf{Y} \mid \mathbf{X})$ generating a forecasted sequence $\mathbf{Y} = [\hat{\mathbf{y}}_1, \dots, \hat{\mathbf{y}}_H]$ conditioned on historical context $\mathbf{X}$. The policy is optimized against an objective combining an empirical or domain-specific reward function $R(\mathbf{X}, \mathbf{Y})$ with a Kullback-Leibler (KL) divergence constraint against a frozen reference foundation model π_{ref} :

$$\max_{\theta} \mathbb{E}_{\mathbf{X} \sim \mathcal{D}, \mathbf{Y} \sim \pi_{\theta}} [R(\mathbf{X}, \mathbf{Y})] - \beta \mathbb{D}_{\text{KL}} (\pi_{\theta}(\cdot \mid \mathbf{X}) \parallel \pi_{\text{ref}}(\cdot \mid \mathbf{X})), \quad (26)$$

where $\beta > 0$ controls the exploration-exploitation trust region, preventing reward collapse and catastrophic forgetting.

Recent landmark works instantiate post-training policy optimization across time-series foundations:

- **TimeRFT: Reinforcement Fine-Tuning for Time Series Foundation Models (Li et al., 2026) [111]:** Li et al. identify that standard SFT on downstream time series suffers from temporal distribution shift between historical training windows and future test horizons. To unlock generalizable adaptation, TimeRFT introduces a two-fold RL fine-tuning strategy: (1) a quality-aware temporal reward mechanism that provides fine-grained, step-wise credit assignment $r_t = -\alpha|y_t - \hat{y}_t| - \gamma \cdot \text{sign}(\Delta y_t \cdot \Delta \hat{y}_t)$, explicitly penalizing directionally incorrect inflection predictions; and (2) a difficulty-aware sample selection policy that prioritizes non-stationary training windows exhibiting high spectral entropy. Evaluated across major TSFM backbones (TimesFM, Chronos, UniTS), TimeRFT achieves up to 18.6% lower forecasting error than standard full fine-tuning and LoRA under severe out-of-distribution regime shifts.
- **TimeHF: Billion-Scale Foundation Models Guided by Human Feedback (Qi et al., 2025) [112]:** Qi et al. pioneer the first 6-billion-parameter foundation model aligned via human feedback for enterprise operations. Recognizing that statistical loss metrics (MSE, MAE) fail to capture domain experts’ tacit operational knowledge, TimeHF introduces Timeseries Policy Optimization (TPO). TPO trains a preference reward model on pairwise expert comparisons

ranking alternative forecast trajectories under asymmetric replenishment penalties (holding cost vs. lost-sales cost). Applied to automated supply chain replenishment across 20,000+ commercial products, TimeHF demonstrates a 33.21% improvement in decision accuracy over deep baselines, proving that human preference alignment fundamentally reshapes commercial forecasting.

- **COUNTS**: Eliciting Chain-of-Thought Reasoning via Reinforcement Learning (Parker et al., 2025) [113]: Parker et al. address the challenge of eliciting explicit, verifiable Chain-of-Thought (CoT) reasoning for time series in LLMs. Traditional LLMs applied to numerical sequences output blunt predictions without intermediate rationales. COUNTS tokenizes continuous waveforms into discrete high-fidelity symbols using a Residual Vector-Quantized VAE (RVQ-VAE). Subsequently, COUNTS optimizes the model using Group Relative Policy Optimization (GRPO) with verifiable outcome rewards on analytical problems (e.g., ECG-QA arrhythmia detection, UCR classification, Context-Is-Key forecasting). By rewarding correct mathematical reasoning traces prior to final answers, COUNTS dramatically outperforms standard prompting, proving that RL can instill deliberate, verifiable multi-step cognitive reasoning over temporal signals.
- **TimeMaster**: Training Time-Series Multimodal LLMs to Reason via RL (Zhang et al., 2025) [114]: Zhang et al. formulate TimeMaster, leveraging RL to train multimodal large language models (such as Qwen2.5-VL) to reason over visual line plots and multi-channel physiological sensor signals (EMG, ECG, HAR). By defining step-by-step diagnostic rewards, TimeMaster eliminates visual hallucinations and aligns textual diagnoses with underlying physiological waveforms.
- **LAST SToP**: Stochastic Soft Prompting for Asynchronous Sequences (Gupta et al., 2025) [115]: Gupta et al. bridge LLM reasoning with irregular time-stamped event sequences in the LAST SToP framework. Framing irregular sequences as Language-modeled Asynchronous Time Series (LASTS), the method introduces Stochastic Soft Prompting (StoP), a lightweight prompt optimization policy that dynamically samples prefix embeddings from learned distributions. StoP significantly outperforms QLoRA on irregular clinical ICU trajectories and sensor failure benchmarks while preserving the base LLM’s pre-trained world knowledge.

6 Evaluations, Benchmarks, and Critical Perspectives

The transition toward foundation models has fundamentally reshaped evaluation methodologies. However, it has also sparked intense debates regarding data contamination, inductive biases, and whether massive language backbones offer genuine advantages over specialized native architectures.

6.1 Standardized Foundation Benchmarks

Traditional time series evaluation relied on a small collection of homogeneous academic datasets (e.g., ETT, Weather, Electricity, Traffic). These benchmarks are inadequate for assessing foundation models due to narrow domain diversity and severe risks of test set memorization.

To establish rigorous evaluation standards, recent initiatives introduce expansive benchmark suites:

- **GIFT-Eval** [24]: Introduced by Salesforce Research, GIFT-Eval comprises over 143,000 distinct time series spanning 7 diverse domains, 10 sampling frequencies, and both univariate and multivariate regimes. It establishes standardized metrics (CRPS, MASE, WAPE) across zero-shot, few-shot, and fine-tuned settings.
- **fev-bench** [116]: Introduced by Shchur et al. (AWS and AutoGluon), fev-bench formulates a realistic benchmark encompassing 100 forecasting tasks across 7 domains. Crucially, 46 of these tasks incorporate dynamic covariates (e.g., promotional flags, weather drivers, calendar events) evaluated via the open-source fev library using bootstrapped confidence intervals and win-rate skill scores, directly addressing the historic deficiency where foundation models underperformed simple baselines by omitting exogenous features.
- **It’s TIME** [117]: Qiao et al. (joint work between Tsinghua THUML and Ming Jin’s group) establish next-generation benchmarks addressing multi-granularity task evaluation and rigorous train-test contamination audits, establishing standardized leakage-free splits across diverse operational horizons.
- **Beyond Numerical Time Series** [118]: Chen et al. establish a unified multimodal evaluation framework pairing continuous quantitative sensors with heterogeneous textual descriptions, news contexts, and structural calendar meta-data.
- **AION** [119]: Zhan, Jin et al. formulate next-generation benchmarks moving beyond static univariate forecasting to evaluate agentic tool-use, multi-step structured reasoning, and contextual decision support over temporal observations.
- **TimesX** [120]: Published at ICML 2026 by researchers from Georgia Tech and Google, TimesX provides a standardized benchmark for context-enriched multimodal forecasting, pairing high-frequency sensors with text narratives, calendar events, and exogenous market drivers across diverse global domains.
- **TimeSeriesExam** [121]: Developed by the CMU MOMENT team, TimeSeriesExam formulates a diagnostic capability examination for foundation models. Rather than evaluating generic mean-squared errors on arbitrary test splits, it probes specific foundational abilities: trend extrapolation, seasonality decomposition, phase-shift invariance, and anomaly localization.
- **LiveHouse-TS** [122]: An open-world living benchmark platform that streams fresh real-world telemetry in real time, preventing static benchmark memorization and providing continuous evaluation under non-stationary concept drift.
- **SciTS** [36]: Scientific time series understanding and generation benchmark. Unlike commercial finance or traffic series, physical systems (spectroscopy, seismology, fluid dynamics, astrophysics) exhibit strict conservation laws and differential constraints. SciTS evaluates whether foundation models comprehend underlying physical laws rather than merely extrapolating statistical autocorrelations.
- **Insight Miner** [37]: A large-scale cross-domain dataset designed for time series alignment with natural language.

Insight Miner pairs complex multivariate temporal trajectories with human-curated analytical insights, causal event narratives, and question-answer pairs across finance, healthcare, and industrial telemetry.

- Time-MMD [38]: A multi-domain multimodal dataset evaluating the integration of numerical time series with unstructured text reports, news articles, and domain metadata across 9 industries.
- Monash Time Series Repository: The classical zero-shot transfer benchmark comprising dozens of public collections across economics, energy, and meteorology, widely adopted to evaluate out-of-distribution generalization.

6.2 Empirical Zero-Shot Benchmark Comparison

To resolve conflicting performance claims in the literature, Table 4 synthesizes reported zero-shot performance across representative foundation models and LLM adaptations on GIFT-Eval, fev-bench, and Monash.

Inspection of Table 4 yields four foundational insights:

- 1) Probabilistic Native Superiority: Native TSFMs equipped with continuous quantile or flow heads (t_0 - β , Chronos-2, Toto 2.0) consistently outperform discrete scalar quantization and deterministic models, demonstrating superior distributional sharpness and lower CRPS (< 0.48 on GIFT-Eval).
- 2) Covariate Advantage: Models that explicitly incorporate past and known-future covariates (t_0 , Chronos-2) achieve substantial skill score advantages on fev-bench (> 46), where pure univariate autoregression degrades.
- 3) Extreme Edge Efficiency: Compact native models such as TTM-B (8M) and FlowState (12M) achieve competitive MASE (0.71–0.74) while operating at 0.1% the parameter footprint of 7B reprogrammed language models.
- 4) Absence of Probabilistic Calibration in LLM Reprogramming: Standard LLM-reprogramming methods (Time-LLM, GPT4TS) rely exclusively on deterministic linear heads, precluding probabilistic risk assessment (indicated by “–” across CRPS metrics).

6.3 Critical Perspectives: Are LLMs Actually Useful for Time Series?

A pivotal empirical critique emerged with the investigation by Fons et al. [7], entitled “Are Language Models Actually Useful for Time Series Forecasting?” Through extensive ablation experiments across multiple LLM-based forecasters (including GPT4TS [20] and Time-LLM [19]), the authors isolated the impact of pre-trained language weights by comparing:

- 1) Frozen pre-trained language weights (e.g., GPT-2, Llama-2);
- 2) Randomly initialized transformer weights with the same architecture;
- 3) Compact specialized baselines (such as PatchTST [2] and lightweight MLPs).

Their empirical findings revealed striking conclusions:

- Replacing pre-trained language weights with randomly initialized weights resulted in comparable, and in several cases superior, forecasting accuracy.
- Much of the empirical gains reported in early LLM4TS literature originated from the structural benefits of subseries

patching and normalization, rather than any semantic or temporal reasoning inherent in linguistic pre-training.

- In terms of operational efficiency, frozen 7B parameter language models require up to $100\times$ more memory and $50\times$ longer inference latency than compact native models like TTM [13] (8M parameters) or Timer [15] (84M parameters), which consistently achieve competitive or superior accuracy.

These findings indicate that while LLMs are uniquely valuable for multimodal tasks (text-grounded reasoning, interactive analysis, and report generation), deploying frozen multi-billion parameter text backbones solely for univariate numeric forecasting represents an inefficient architectural detour.

Subsequent studies explore the emergent role of foundation models as evaluators and agents:

- TimeVista [123]: Chen et al. (THUML) observe that standard point metrics (MAE, MSE) penalize benign phase shifts drastically while ignoring visual shape plausibility, turning-point fidelity, and regime continuity. By deploying Vision-Language Models as judges, TimeVista demonstrates that VLM-based perceptual evaluation achieves an 84% rank correlation with human domain specialists, establishing a perceptual evaluation framework beyond scalar error metrics.
- Forecast Workflow Bench [124]: Nagashima investigates whether LLM agents can orchestrate time-series foundation models under realistic operational tool budgets. Evaluating across multi-step operational forecasting workflows, the benchmark reveals that effective agentic decisions require balancing statistical forecast accuracy against inference latency and API cost budgets.

6.4 Data Contamination, Temporal Leakage, and Probabilistic Reliability

6.4.1 Benchmark Contamination Audits

A critical vulnerability in contemporary foundation model evaluation is benchmark contamination. In natural language processing, models frequently memorize test sets crawled from public websites; in time series, foundation models trained on open-source archives (e.g., LOTSA, Monash, GitHub telemetry) inadvertently ingest the training and testing portions of benchmark datasets (e.g., ETT, Weather) [25].

Gong et al. [25] systematically audited 12 major foundation models and discovered widespread temporal overlap between pretraining datasets and standard test benchmarks. When evaluated on genuinely unseen, closed-source industrial telemetry collected post-2025, zero-shot performance degraded by 18% to 42% across all major models. This underscores the urgent need for living, dynamic benchmarks (such as LiveHouse-TS [122]) where evaluation datasets are continuously refreshed.

6.4.2 Probabilistic Reliability and Quantile Calibration

Beyond mean error metrics, real-world deployment in safety-critical sectors (energy grids, algorithmic trading, intensive care) demands calibrated uncertainty intervals. In a recent rigorous audit, Li et al. [26] investigated the probabilistic reliability of zero-shot TSFMs. Their findings reveal a concerning dichotomy: models achieving competitive point forecast errors

TABLE 4
Zero-Shot Empirical Benchmark Comparison Across Foundation Models and LLM Baselines.

Model	Backbone / Family	Paradigm	Params	GIFT-Eval		fev-bench Skill ↑	Monash Repository			Uncertainty Head
				CRPS ↓	MASE ↓		MASE ↓	WAPE ↓	CRPS ↓	
Chronos-Large [3]	T5-Large	Native TSFM	710M	0.489	0.732	—	0.812	0.449	0.380	Categorical Bins
Chronos-2 [8]	T5-Base Hybrid	Native TSFM	120M	0.468	0.672	48.2	0.795	0.435	0.362	Continuous Hybrid
TimesFM-200M [4]	Custom Decoder	Native TSFM	200M	0.518	0.742	—	0.820	0.455	—	Quantile Head
MOIRAI-Large [5]	Custom Enc-Dec	Native TSFM	311M	0.505	0.724	43.5	0.834	0.461	0.412	Parametric Mixture
Lag-Llama [14]	LLaMA Decoder	Native TSFM	2.4M	—	—	—	0.942	0.510	0.521	Student's t
MOMENT-Large [12]	T5 Encoder	Native TSFM	385M	—	0.768	—	0.871	0.485	—	Deterministic Recon
Timer [15]	GPT-2 Custom	Native TSFM	84M	—	0.738	—	0.835	0.465	—	Deterministic Patch
Sundial [17]	Multi-Rate Dec	Native TSFM	1.5B	0.482	0.698	—	0.805	0.442	0.375	Continuous Flow
t_0 - β [30]	Axial Transformer	Native TSFM	256M	0.4738	0.6865	46.7	0.801	0.438	0.369	Quantile Head
FlowState [41]	Neural ODE / Flow	Native TSFM	12M	0.491	0.712	—	0.828	0.456	0.388	Continuous Flow
Toto 2.0 [45]	Wavelet Decoder	Native TSFM	1.2B	0.465	0.668	49.0	0.788	0.430	0.355	Quantile Head
TTM-B [13]	TSMixer Blocks	Native TSFM	8M	—	0.745	—	0.840	0.468	—	Deterministic Point
Time-LLM [19]	LLaMA-7B Reprog	LLM4TS	7B	—	0.785	—	0.885	0.482	—	Deterministic Linear
GPT4TS / OFA [20]	GPT-2 Reprog	LLM4TS	124M	—	0.798	—	0.892	0.491	—	Deterministic Linear
PatchTST [2]	Channel-Indep Trans	Supervised Baseline	1.2M	—	0.810	—	0.908	0.502	—	Deterministic Linear

Note: Performance figures reflect published zero-shot evaluation protocols on official benchmark suites. Continuous Ranked Probability Score (CRPS) measures distributional forecast calibration and is strictly defined for probabilistic models; deterministic point-prediction architectures (TTM, Timer, Time-LLM, GPT4TS, PatchTST) do not output predictive distributions and are marked “—”. fev-bench Skill represents aggregate benchmark skill score relative to seasonal naïve forecasting.




Fig. 2. Genealogy and chronological release timeline of representative time series foundation models and repurposed language model frameworks from 2022 to 2026.

often display severe quantile miscalibration. Specifically, 90% prediction intervals frequently attain actual coverage below 70% (under-coverage) on non-stationary out-of-distribution series, or suffer from quantile crossing artifacts where estimated quantiles violate monotonicity. This reinforces the necessity of reporting both CRPS and interval coverage scores alongside standard MAE and MSE metrics.

6.5 Empirical Scaling Laws and Meta-Regression in Temporal Data

Do foundation models for time series exhibit the predictable power-law scaling observed in NLP? Recent empirical studies on large pre-trained models—notably Time-MoE [40] (2.4B parameters, 300B points), Sundial [17] (1.5B parameters), Timer-S1 [18] (8.3B parameters), and Toto 2.0 [45] (2.5B parameters)—demonstrate that when data quality and volume are rigorously maintained, validation loss follows a

bivariate power law as a joint function of non-embedding parameter count N and unique training tokens D :

$$L(N, D) = \left(\frac{N_c}{N}\right)^{\alpha_N} + \left(\frac{D_c}{D}\right)^{\alpha_D} + L_0, \quad (27)$$

where N_c, D_c denote characteristic scale constants and L_0 represents the irreducible entropy of stochastic observation noise. Given the training compute budget $C \approx 6ND$, the optimal parameter allocation follows $N_{\text{opt}} \propto C^{\frac{\alpha_D}{\alpha_N + \alpha_D}}$ and $D_{\text{opt}} \propto C^{\frac{\alpha_N}{\alpha_N + \alpha_D}}$.

Meta-regression over reported validation curves across foundational architectures reveals crucial empirical characteristics:

- Fitted Exponents: Time-MoE [40] reports $\alpha_N \approx 0.089$ and $\alpha_D \approx 0.112$ across models ranging from 50M to 2.4B parameters on Time-300B; Sundial [17] reports $\alpha_N \approx 0.076$ on UTSD-3; and Toto 2.0 [45] reports $\alpha_N \approx 0.081$ and $\alpha_D \approx 0.098$ under Maximal Update Parametrization




Fig. 3. Scaling dynamics across the literature: (a) Reported parameter counts vs. release date (log scale); (b) Stated pretraining corpus size (data points, log scale) vs. release date.




Fig. 4. Meta-review statistics: (a) Publications per year broken down by model paradigm (stacked bar chart); (b) Distribution of included studies across premier peer-reviewed venues.

(μP). In comparison, LLM language pretraining exhibits $\alpha_N \approx 0.076\text{--}0.088$ and $\alpha_D \approx 0.095\text{--}0.120$. This confirms that temporal transformers obey bona fide neural scaling laws.

- The Temporal Entropy Saturation Barrier: In time series, α_N exhibits early saturation when scaling beyond 1B parameters if trained solely on public sensor repositories. Because real-world telemetry frequently contains repetitive stationary waveforms with low Kolmogorov complexity, models quickly exhaust novel temporal structures. To sustain power-law scaling ($\alpha_N > 0$), state-of-the-art pipelines (Chronos, Time-MoE, Tabby [50]) inject synthetic generative processes (mixtures of non-linear differential equations and Gaussian process kernels), proving that synthetic diversity is mandatory to maintain compute-optimal scaling in time series.

6.6 Computational Profiling, Operational Complexity, and Real-Time Inference

While accuracy metrics (CRPS, MASE) dominate research papers, practical adoption across industrial SCADA telemetry, financial high-frequency execution, smart energy grids, and resource-constrained edge microcontrollers is governed fundamentally by computational efficiency, memory footprint, and inference latency. Table 5 formalizes the algorithmic complexity, memory overhead, empirical latency, and operational deployment envelopes across the six foundational time-series paradigms.

Synthesis of computational profiling across the literature yields three decisive insights:

- 1) The Autoregressive Generation Bottleneck: Models relying on scalar or token-wise autoregression (e.g., original Chronos [3] iterating over 4096 bins) incur H sequential attention forward passes to forecast an H -step horizon. In production benchmarks, Chronos-Bolt resolves this bottle-

TABLE 5
Operational Complexity, Inference Latency, and Memory Footprints Across Foundation Paradigms.

Foundational Paradigm	Exemplar Models	Time Complexity (Inference)	KV-Cache / Peak Memory	Latency (ms)	Generation Mechanism	Hardware / Deployment Envelope
Autoregressive Token / Point Transformer	Chronos-Large [3], AutoTimes [46]	$\mathcal{O}(H \cdot L \cdot d + H^2 \cdot d)$	High ($\mathcal{O}((L+H) \cdot d)$ KV-cache)	120–450	Step-by-step sequential sampling (H passes)	Server (A100/H100); GPU latency-prohibitive for sub-second control
Patch-Based Encoder-Decoder	Chronos-Bolt, MOIRAI TimeXer [53]	$\mathcal{O}((L/P)^2 \cdot d + (H/P) \cdot (L/P) \cdot d)$	Medium ($\mathcal{O}((L/P) \cdot d)$ patch cache)	8–25	Direct multi-patch projection (1 or few passes)	Workstation GPU / Edge GPU (Jetson Orin); 250× faster than vanilla Chronos
Multi-Scale Channel Mixers	TTM, LightGTS, TimeMixer++ [44]	$\mathcal{O}(L \cdot d)$ (Linear mixing)	Minimal (≤ 50 MB, zero KV-cache)	1.5–5	Single-pass linear time/channel projection	Microcontrollers, IoT gateways, embedded CPU; sub-5ms real-time
Continuous Flow-Matching & Diffusion	FlowState, FLAME, UrbanDiT [60]	$\mathcal{O}(K \cdot (L/P)^2 \cdot d)$ (K ODE steps, $K \in [10, 50]$)	Medium ($\mathcal{O}((L/P) \cdot d)$, zero KV-cache)	35–120	Multi-step numerical ODE/SDE integration	Standard GPU; viable for batch forecasting, unsuited for sub-10ms loops
Prior-Data Fitted Bayesian Networks	TabPFN-TS, ForecastPFN, ICTSP [52]	$\mathcal{O}((N_{ctx} + H)^2 \cdot d)$ (Context attention)	Medium-High ($\mathcal{O}(N_{ctx} + H)^2$ prompt memory)	15–60	Single-pass exact Bayesian posterior	Workstation GPU / High-RAM CPU; instant edge zero-shot cold start
Reprogrammed Multi-Billion LLMs	Time-LLM [19], OFA [20], LLMTime [6]	$\mathcal{O}(L_{text}^2 \cdot D_{LLM} + H \cdot D_{LLM})$	Extreme (14–80GB VRAM, multi-head KV cache)	500–3500	Text tokenization + multi-layer LLM forwards	Multi-GPU cloud clusters only; completely prohibited on edge devices

neck by shifting from point-wise autoregression to patch-level direct quantile regression, achieving a documented $250\times$ inference speedup and $20\times$ memory reduction while matching or exceeding zero-shot accuracy.

- 2) The Extreme Throughput Gap (Mixers vs. Reprogrammed LLMs): Lightweight architectural designs—exemplified by IBM’s Tiny Time Mixers (TTM [13], 1M–8M parameters) and LightGTS [55] (5.8M parameters)—execute full batch inference in under 5ms on commodity hardware with zero KV-cache memory overhead. In contrast, reprogrammed LLMs (Time-LLM [19], AutoTimes [46]) require multi-second latency per inference batch and allocate over 14GB of GPU VRAM. For high-frequency industrial automation where millisecond response times are mandatory, compact native mixers represent the only viable foundation paradigm.
- 3) In-Context Bayesian Inference vs. Score-Based Diffusion: Prior-Data Fitted Networks (TabPFN-TS [51]) offer unique operational trade-offs: while context-sequence attention scales quadratically with history length, they execute zero gradient updates, eliminating all fine-tuning optimizer overhead. Conversely, continuous score-based flow matching (FlowState [41], UrbanDiT [60]) demands multiple numerical ODE integration steps ($K \in [10, 50]$), exchanging inference latency for mathematically continuous, sampling-rate-equivariant trajectories.

6.7 Energy Consumption, Integer Quantization Degradation, and Edge Deployment Profiles

While academic benchmarks focus overwhelmingly on predictive error (MSE, CRPS), the viability of deploying foundation models in cyber-physical systems (CPS), Internet-of-Things (IoT), and embedded industrial controllers is strictly determined by energy consumption (Joules per inference), silicon area constraints, and numerical quantization resilience [125], [126].

To systematically address cross-platform hardware discrepancies, recent literature establishes three pivotal advancements:

- 1) HoliBench: Holistic Cross-Platform Profiling: Chakrabarti et al. [125] introduce HoliBench, a standardized benchmarking and deployment toolkit that jointly characterizes accuracy, latency, and energy across seven device classes, eight inference backends, and three quantization levels

(FP16, INT8, INT4). HoliBench uncovers that integer post-training quantization (PTQ) reduces latency only on architectures equipped with specialized low-precision compute pipelines (e.g., INT8 Tensor Cores or NEON vector units); on general-purpose microcontrollers lacking integer SIMD, quantization merely reduces memory footprint while potentially increasing execution latency due to dequantization overhead. Furthermore, HoliBench demonstrates that standalone single-model profiles compose with high fidelity (1.2% latency and 2.5% power error) under co-resident execution, enabling optimal multi-model pipeline configuration without combinatorial profiling.

- 2) Resource-Aware Mixed-Precision FPGA Quantization: Ling et al. [126] resolve the hardware-level bottleneck of deploying time-series Transformers onto resource-constrained embedded FPGAs (Xilinx Spartan-7 XC7S15). By formulating a resource-aware mixed-precision quantization methodology (allocating INT8 to sensitive cross-attention maps and INT4 to feed-forward blocks) and optimizing block RAM (BRAM) storage templates, they achieve hardware resource estimation within 3% of physical silicon synthesis, unlocking deployability for previously non-synthesizable transformer configurations.
- 3) Carbon-Aware Dynamic Dispatch (FM-CAC): Yang et al. [127] formulate FM-CAC, establishing a proactive carbon-aware control framework for battery-buffered edge AI systems. By deploying edge-friendly zero-shot TSFMs to forecast temporal marginal emissions and grid carbon intensity, FM-CAC feeds predictions into a dynamic programming solver with deferred cost attribution. This decouples energy harvesting from compute execution, reducing carbon emissions by up to 65.6% while maintaining near-maximum inference accuracy across continuous edge vision and time-series pipelines.

As detailed in Table 6, lightweight architectures (e.g., Cheraghinia MLP [65], TTM [13], and APEX-Edge [64]) consume between 0.18 mJ and 28.5 mJ per inference, operating comfortably within battery-buffered and energy-harvesting IoT regimes with negligible ($< 1.2\%$) degradation under INT8 quantization. In contrast, multi-billion parameter LLMs demand over 18 Joules per forward pass and over 14 GB of VRAM, underscoring that compact, quantization-resilient native models are mandatory for sustainable edge intelligence.

TABLE 6
Empirical Energy Consumption, Quantization Degradation, and Edge Hardware Deployment Profiles across Foundation Models.

Model & Paradigm	Parameters	Precision Format	FLOPs / Inference	Energy / Inference	Peak RAM	Accuracy Degradation (Δ MSE)	Target Hardware Deployment Envelope
Cheraghinia MLP [65]	21K	FP32 / INT8	4.2×10^4	0.18 mJ	0.12 MB	< 0.3% (INT8 PTQ)	ARM Cortex-M4/M7 Micro-controller, Nordic nRF52840; sub-millisecond edge IoT
TTM-1M (IBM) [13]	1.0M	FP16 / INT8	2.1×10^6	3.4 mJ	4.2 MB	+0.8% (INT8 PTQ)	Embedded Linux gateways, Raspberry Pi Zero 2W; real-time telemetry
APEX-Edge [64]	10.5M	FP16 / INT8	2.2×10^7	28.5 mJ	24.8 MB	+1.2% (INT8 PTQ)	Enterprise Wi-Fi APs, Cisco Catalyst gateways; sub-second network anomaly audit
Tiny-TSM [63]	23.0M	FP16 / INT8	4.8×10^7	52.0 mJ	49.0 MB	+1.5% (INT8 PTQ)	Single-board computers (Raspberry Pi 4/5, Coral Dev Board); edge forecasting
Mixed-Precision Transformer [126]	1.8M	INT8 / INT4	3.6×10^6	4.1 mJ	1.8 MB (BRAM)	+1.9% (Mixed INT8/INT4)	Embedded FPGA (Xilinx Spartan-7 XC7S15); strict hardware-timed SCADA loops
Chronos-Bolt Small	46.0M	FP16 / INT8	9.5×10^7	112 mJ	98 MB	+1.4% (INT8 PTQ)	Edge AI accelerators (NVIDIA Jetson Orin Nano, Google Coral); high-throughput
MOIRAI-Small [5]	14.0M	FP16 / INT8	3.1×10^7	74 mJ	62 MB	+1.1% (INT8 PTQ)	Industrial IPCs, Jetson Xavier NX; any-variate sensor fusion
TimesFM-200M [4]	200M	FP16 / INT8	4.2×10^8	580 mJ	480 MB	+2.1% (INT8 PTQ)	Edge servers, industrial workstations; multi-horizon continuous quantile forecasting
Time-LLM (7B) [19]	7.0B	FP16 / INT4	1.5×10^{10}	18.4 J	14.2 GB	+3.8% (INT4 AWQ)	Cloud multi-GPU servers (A100/H100); strictly prohibited on edge/battery devices

6.8 Causal Auditing, Test-Time Adaptation Benchmarks, and Living Leaderboards

As foundation models increasingly transition from academic prototypes to mission-critical infrastructure, evaluation methodologies must evolve beyond aggregate static point-error metrics (MSE, MASE) toward verifying causal structural fidelity, resilience under test-time drift, and continuous contamination-resistant leaderboards [76], [128]. Table 7 systematically categorizes the core mechanisms, parameter footprints, and capability envelopes across the emerging landscape of test-time adaptation and causal foundation modeling.

Synthesis of these evaluative frontiers yields four critical insights:

- 1) Causal Audits and the Persistence Bias Failure Mode: Jander, van Heeswijk, and Mes [76] demonstrate that transitioning from bespoke models to foundation forecasters creates systemic concentration risk. By probing state-of-the-art foundation models (Chronos-2 [8], TimesFM-2.5 [4], MOIRAI [5]) using synthetic interventional data generators, their causal audit reveals that foundation models exhibit a severe persistence bias. Specifically, when probed with processes exhibiting delayed interventional responses, sudden regime shifts, or unobserved confounders, models default to extrapolating local lag trends, severely underestimating the magnitude and velocity of structural inflections.
- 2) Realistic Benchmarking for Temporal Causal Discovery: Cheng et al. [129] establish CausalTime (NeurIPS 2023), resolving a core bottleneck in temporal causal validation. Prior benchmarks relied on oversimplified linear autoregressive models with Gaussian noise, producing synthetic evaluations disconnected from real-world physics. CausalTime utilizes deep invertible generative neural networks and normalizing flows to capture complex nonlinear dynamics and heteroskedasticity while deriving exact

ground-truth causal DAGs, creating a rigorous standard benchmark for evaluating causal graph discovery.

- 3) Principled Test-Time Adaptation Benchmarking (FAC): Wang et al. [128] critique existing TTA evaluations for subtle future-information leakage and ad-hoc batch scheduling. They establish a standardized leak-free streaming benchmark suite and introduce Frequency-Aware Calibration (FAC). FAC proves that adapting models directly in the raw temporal domain frequently amplifies high-frequency measurement noise into destructive gradient shocks; conversely, isolating adapter updates to orthogonal low- and mid-frequency Fourier sub-bands preserves base model stability while yielding consistent accuracy improvements across non-stationary streaming horizons.
- 4) Living Leaderboard Ecosystems (The TIME Benchmark): Addressing the rapid obsolescence of static benchmark datasets, Qiao et al. [117] introduce the TIME benchmark (50 fresh datasets, 98 tasks) and the public Hugging Face TIME-Leaderboard (<https://huggingface.co/spaces/Real-TSF/TIME-leaderboard>). By pairing automated weekly evaluation pipelines with pattern-level time series decomposition, TIME evaluates foundation models across multi-scale periodicity, non-stationarity, and noise regimes, establishing a living, contamination-resistant community reference standard for zero-shot time-series intelligence.

6.9 Living Streaming Benchmarks, Temporal Generalization, and Agentic Multi-Turn Evaluation

Static benchmark splits (e.g., standard 70/10/20 train-val-test partitioning on historical repositories) have increasingly become liabilities in foundation model evaluation. Beyond the acute risk of pretraining data contamination and lookahead leakage, static benchmarks fail to measure how models perform under continuous open-world temporal evolution [130].

TABLE 7
Systematic Synthesis of Test-Time Adaptation (TTA) and Causal Time-Series Foundation Models.

Framework	Key Reference	Core Mechanism & Adaptation / Causal Objective	Base Model	Param Footprint	Inference Overhead	Target Shift / Causal Capabilities
Category A: Test-Time Adaptation (TTA) under Non-Stationary Drift						
TSF-TTA	Kim et al. [73] (AAAI'25)	Online update of normalization statistics and linear projection heads via temporal consistency loss $\mathcal{L}_{TTA} = \ \hat{\mathbf{y}}_t - \mathcal{T}(\hat{\mathbf{y}}_{t-1})\ _2^2$	Frozen PatchTST / DLinear	Small (< 5M)	+12% latency	Abrupt mean/variance non-stationarity without source data
AdaNODEs	Dang et al. [74] (ICASSP'26)	Continuous latent trajectory adaptation via Neural ODEs: $\frac{d\mathbf{z}(t)}{dt} = f_\theta(\mathbf{z}(t), t)$ with self-supervised reconstruction	Any Frozen Forecaster	Minimal (0.15M)	+8% latency	Source-free continuous non-linear temporal drift
RG-TTA	Kumar et al. [75] (arXiv'26)	Regime-guided meta-control dynamically detecting stationary vs. non-stationary regimes, modulating gradient steps $\eta_t = \sigma(\text{RegimeGate}(X_t))$	Pretrained Transformer	Zero extra params	+5% latency	Multi-regime streaming series; prevents catastrophic forgetting
FAC Framework	Wang et al. [128] (arXiv'26)	Frequency-Aware Calibration projecting updates onto orthogonal spectral sub-bands to suppress high-frequency noise explosion	Pretrained TSFM	Lightweight	+10% latency	Leakage-free streaming adaptation; eliminates noisy gradient divergence
Category B: Causal Discovery, Structural Priors, and Counterfactual Foundation Models						
Causal-PT	Stein et al. [77] (arXiv'24)	Direct transformer mapping from multivariate temporal tokens to causal DAG adjacency matrix $\mathbf{A} \in \{0, 1\}^{D \times D}$ pre-trained across 10M SCMs	Transformer Encoder	25M	Single forward pass	Zero-shot discovery of directed causal graphs from observational series
CausalTimePrior	Thumm & Chen [78] (arXiv'26)	Prior-Data Fitted Network (PFN) pre-trained on 500M synthetic interventional TSCMs to approximate exact in-context posterior $P(\mathbf{Y} do(\mathbf{X}), \mathcal{D}_{obs})$	Transformer PFN	85M	Single forward pass	In-context causal effect estimation and zero-shot counterfactual forecasting
CaTSG	Xia et al. [79] (arXiv'25)	Structural causal score-based diffusion model conditioning reverse denoising trajectories on time-varying causal graphs $\mathcal{G}_t$	Causal Diffusion Net	45M	Multi-step ODE/SDE	Controllable trajectory generation under interventions ($do(X_t = \alpha)$)
CausalTime Bench	Cheng et al. [129] (NeurIPS'23)	Deep normalizing flow generator providing realistic multivariate time series paired with verifiable ground-truth causal graphs	Benchmark Generator	N/A	Offline evaluation	Ground-truth benchmark suite for time-series causal discovery algorithms
Jander Audit	Jander et al. [76] (arXiv'26)	Synthetic interventional audit harness probing persistence bias and structural distortion across foundation models	Chronos-2, TimesFM, MOIRAI	N/A	Offline audit	Uncovers persistence bias and regime switch degradation in TSFMs

Simultaneously, standard single-step error metrics (MSE, MAE, CRPS) cannot evaluate multi-turn agentic workflows where conclusions emerge from cumulative dialogue turns, tool executions, and state reflections [100].

Table 8 provides a systematic taxonomy and architectural comparison across the emerging generation of agentic time-series reasoning frameworks and living temporal benchmarks.

These living and agentic evaluation frameworks illuminate two fundamental paradigms:

- 1) Continuous Live Evaluation vs. Static Overfitting: Garza et al. [130] instantiate Impermanent on top-400 GitHub activity streams. Because software repositories experience organic shocks (major version releases, viral social adoption, tooling migrations), they exhibit intense non-stationarity without synthetic artifacts. Crucially, Impermanent scores forecasts sequentially over time on live data streams with daily leaderboard updates (<https://impermanent.timecopilot.dev>). Their findings expose that zero-shot foundation models showing near-optimal static error rates frequently deteriorate when evaluated under sustained streaming evolution, underscoring that static benchmarks reward memorization and lookahead leakage rather than true temporal generalizability.
- 2) Multi-Turn Dialogue and Decision Gaps: Kong et al. [100] construct TimeSage-MT, bridging the gap between single-step predictions and practical decision-making workflows. Across 2,680 dialogue turns in eight domains (finance, energy, transportation, healthcare, meteorology, manufac-

turing, retail, and server telemetry), they discover that even frontier closed-source LLMs (GPT-4o, Claude 3.5 Sonnet) suffer sharp performance drops exceeding 35% when progressing from exploratory questions to complex multi-step decision tasks. These failures stem primarily from three bottlenecks: loss of intermediate memory over long dialogue horizons, inability to calibrate forecast uncertainty across tool invocations, and domain hallucination when resolving conflicting statistical evidence.

6.10 Data Contamination, Pretraining Familiarity Bias, and Automated Continuous Auditing

The rapid proliferation of time-series foundation models has exposed an acute methodological challenge: benchmark data contamination and temporal leakage. Unlike natural language corpora where test sets can be filtered via string matching, time series consist of continuous, normalized numerical sequences that undergo varying degrees of preprocessing (standardization, differencing, patch discretization). Furthermore, because the pretraining datasets of many prominent models (e.g., UTSD, Chronos-Data, Monash) aggregate open-source archives, classic benchmarks (ETT, Weather, Electricity, Traffic) are frequently ingested during pretraining, rendering reported zero-shot performance suspect.

Recent literature systematically tackles this vulnerability through automated auditing frameworks, leakage-free simulation gyms, and temporal familiarity metrics:

TABLE 8
Taxonomy and Architectural Profiling of Agentic Time-Series Reasoning Frameworks and Living Benchmarks.

Framework	Key Reference	Core Role & Paradigm	Base Model / Backbone	Toolset & Action Space	State & Memory Mechanism	Evaluated Benchmarks & Capabilities
Category A: Interactive Reasoning and Tool-Augmented Agents						
TimeInteract	Pan et al. [101] (arXiv'26)	Streaming Real-Time Interactive Intelligence	Qwen2-7B Backbone with Dual-View TS Encoder	Streaming perception, silent/response gating, decoupled generation	Dual-view short/long-term rolling perceptual memory	StreamTSI-34K (34.5K episodes); near-zero stall; 2.15× speedup
Cast-R1	Tao et al. [102] (arXiv'26)	Tool-Augmented Sequential Decision Forecaster	Qwen2.5-7B (SFT + Multi-Turn RL)	Seasonal-trend decomposition, stationarity tests, lightweight model calls	Episodic decision memory maintaining historical evidence	Complex non-stationary benchmarks; multi-turn self-reflective forecasting
TimeART	Wu et al. [103] (arXiv'26)	Agentic Data Scientist for Time-Series QA	8B Time Series Reasoning Model (TSRM) on Llama-3	Data science tool suite, statistical tests, plotting, forecasting heads	Trajectory memory with experience replay	TimeToolBench (100k expert trajectories); high-precision TSQA
TS-Reasoner	Ye et al. [104] (arXiv'24)	Domain-Specialized Multi-Step Inference Agent	7B / 14B Language Models	Domain-specific computational routines, symbolic reasoning tools	Iterative error feedback loop with constraint validator	TimeSeriesExam; complex multi-step compositional inference
DCATS	Yeh et al. [105] (arXiv'25)	Data-Centric AutoML Agent for Time Series	GPT-4 / Llama-3-70B Planner	Data cleaning routines, anomaly imputation, pipeline configurators	Pipeline state tracker optimizing data quality	Large-scale traffic volume datasets; 6% average error reduction across models
Category B: Living Temporal Benchmarks and Multi-Turn Diagnostic Suites						
Impermanent	Garza et al. [130] (arXiv'26)	Continuous Live Streaming Generalization Benchmark	TimeCopilot Live Platform (Top 400 GitHub Repos)	Daily automated ingestion of PRs, issues, commits, stargazers	Rolling daily window with sequential scoring	Living leaderboard; evaluates temporal robustness against distribution breaks
TimeSage-MT	Kong et al. [100] (arXiv'26)	Multi-Turn Agentic Reasoning Benchmark Suite	Frontier LLMs + Structured TimeSage Agent	Comprehensive time-series skill library (decomposition, anomaly, stats)	Multi-turn conversational dialogue state tracking	240 tasks, 2,680 turns across 8 domains; audits memory and uncertainty failure

- **TSFMAudit: Automated Temporal Contamination Scanning** (Li et al., 2026) [131]: Li et al. develop TSFMAudit, establishing the first formal auditing suite for detecting pretraining contamination in time-series foundation models. TSFMAudit introduces two complementary statistical tests: (1) Quantile-Bin N-Gram Hashing, which maps continuous floating-point sequences into discrete symbolic tokens to detect verbatim sequence memorization, and (2) Mutual Information Permutation Testing, which quantifies non-linear statistical dependence between candidate pre-training segments and benchmark test sets. Auditing multiple prominent open-weight foundation models, TSFMAudit reveals that up to 18.4% of standard benchmark sequences exhibit statistically significant memorization signatures, leading to artificially deflated zero-shot test errors.
- **Pretraining Familiarity Bias Survives Temporal Hold-Outs** (Moghadas & Ghaderi, 2026) [132]: Moghadas and Ghaderi challenge the pervasive research convention that temporal hold-outs (e.g., training on pre-2020 data and testing on post-2020 data) guarantee an uncontaminated evaluation. They show that because physical data-generating processes (e.g., weather cycles, power grid load curves) exhibit stationary dynamical attractors, a foundation model exposed to historical records retains profound pretraining familiarity. When evaluated on temporally separated hold-outs from familiar domains versus genuine cross-domain datasets, foundation models experience performance gaps of up to 42%, proving that reported zero-shot scores frequently conflate cross-domain generalization with domain-specific familiarity.
- **Empirical Stress-Testing on Volatile Infrastructure** (Pan & Ezzat, 2026) [133]: Pan and Ezzat conduct an exhaustive empirical audit of zero-shot foundation models (Chronos-Bolt, TimesFM, Moirai) on electricity price markets (PJM, EPEX, and Australian NEM). Characterized by extreme price spikes and renewable penetration, electricity markets

violate Gaussian stationarity. Their audit uncovers that univariate foundation models degrade by 28% compared to localized gradient-boosted trees (LightGBM) due to their inability to ingest critical exogenous covariates (temperature, transmission capacity). Furthermore, contamination risk is heightened because public market data is ubiquitously crawled into pretraining corpora without provenance tracking.

- **Forecast-Dojo: Replayable Environments for LLM Forecasting Agents** (Ye et al., 2026) [134]: Ye et al. introduce Forecast-Dojo, a leak-free replayable simulation gym that benchmarks multi-agent forecasting systems on resolved prediction market questions (from Polymarket and Metaculus). Forecast-Dojo deterministically pairs real-world questions with dated, timestamped global news archives, strictly eliminating future lookahead bias. In this controlled environment, agents must gather historical evidence, write Python code, and generate calibrated probabilistic forecasts, providing a rigorous testbed for evaluating genuine temporal reasoning without data contamination.

Table 9 synthesizes the frontier across multi-agent deliberation systems, data contamination auditors, and extreme low-bit edge quantization foundations.

6.11 Distribution-Free Conformal Prediction and Statistical Coverage Guarantees

A profound deficiency in contemporary time series foundation models is their vulnerability to uncalibrated probabilistic uncertainty. While foundation models deploy various heads—such as categorical cross-entropy over quantized bins (Chronos), continuous Student’s t / Gaussian parameters (TimesFM, Moirai), or diffusion score matching—their nominal prediction intervals lack rigorous finite-sample coverage guarantees. Under out-of-distribution drift, non-stationary variance scaling, or heavy-tailed shocks, standard foundation model intervals suffer from catastrophic empirical undercov-

TABLE 9
Taxonomy and Empirical Characterization of Multi-Agent Deliberation Systems, Contamination Scanners, and Edge Quantization Foundations.

Framework	Key Reference	Core Paradigm	Category &	Architectural Mechanism	Target Hardware / Scope	Key Metric / Empirical Finding	
Category A: Multi-Agent Swarm Deliberation and Self-Evolving Systems							
MC-Debate	Trirat et al. [107] (arXiv'26)	Multi-Agent Swarm Deliberation		Cross-modal debate across Numerical, and Textual specialists	Server (7B/14B bones) GPUs Back-	−14.2% reasoning error; eliminates 76% of hallucinated trend reversals	
TimeEvo	Yang et al. [108] (arXiv'26)	Failure-Driven Evolution	Self-	Evolutionary agent autonomously synthesizing Python tools on task failure	Server (8B/70B controllers)	Clusters Con-	Dynamic tool synthesis; continuous improvement on TimeQA without human engineering
Traceable-Agent	Kang et al. [106] (arXiv'26)	Traceable Multi-Agent	Enterprise	Four sub-agents with immutable DAG linking 10-K filings to forecasts	Enterprise Infrastructure	Cloud	100% decision traceability; verifiable causal citations for regulatory compliance
MetaCaster	Shen et al. [110] (arXiv'26)	Meta-Harness Synthesis	Forecaster	LLM agent architecting compact models (<1M) via dataset fingerprints	Edge / IoT nodes	Low-power	Matches 100M+ TSFM zero-shot accuracy with 20× lower inference compute
Category B: Extreme Low-Bit Quantization and Edge Inference Foundations							
SparseBinary	Gorbett et al. [70] (arXiv'23)	1-Bit Weight former	Trans-	Extreme binary weights ({−1, +1}) with addition-only ALU accumulation	Embedded / ASICs	FPGAs	87% parameter reduction; multiplication-free multivariate attention
Q-DEQ	Yang et al. [71] (arXiv'26)	Quantized Deep Equilibrium Model		Implicit fixed-point solving ($\mathbf{z}^* = f_\theta(\mathbf{z}^*, \mathbf{x})$) under INT4/INT8	Bare-metal MCUs (<100KB SRAM)		0.3M params; 75 KB footprint; zero dynamical stability loss
SGA	Hu et al. [72] (arXiv'26)	State-Guided Autoregressive Adapter	Autore-	Latent state curvature tracking to scale autoregressive rollout variance	TSFM Backbones (15M Adapter)		Calibrated prediction intervals over long horizons without Monte Carlo sampling
Category C: Data Contamination Auditing, Familiarity Bias, and Replayable Gyms							
TSFMAudit	Li et al. [131] (arXiv'26)	Data Contamination Scanner		Quantile n-gram hashing and mutual information permutation testing	Pretraining Corpora (UTSD, Chronos)		Detects up to 18.4% sequence memorization in public open-weight models
Familiarity-Audit	Moghadas et al. [132] (arXiv'26)	Pretraining Familiarity Evaluator		Pretraining familiarity index across pre/post temporal benchmark splits	Benchmark Archives (ETT, Weather)		Proves temporal hold-outs preserve domain attractor bias (up to 42% gap)
Forecast-Dojo	Ye et al. [134] (arXiv'26)	Replayable Agent Simulation Gym		Prediction markets paired with dated news streams in leakage-free gym	Multi-Agent Evaluation		Leakage-free benchmark for multi-agent temporal decision policies

erage (e.g., nominal 90% prediction intervals achieving only 65–75% empirical coverage).

Conformal Prediction (CP) provides an elegant, model-agnostic, distribution-free mathematical framework that converts any point or quantile forecaster into set-valued prediction regions $\mathcal{C}_{1-\alpha}(\mathbf{X}_{n+1})$ guaranteeing exact finite-sample marginal coverage:

$$\mathbb{P}(\mathbf{Y}_{n+1} \in \mathcal{C}_{1-\alpha}(\mathbf{X}_{n+1})) \geq 1 - \alpha, \quad (28)$$

for a user-specified significance level $\alpha \in (0, 1)$, without requiring Gaussianity, linearity, or asymptotic sample sizes.

Recent breakthroughs systematically integrate conformal inference with time series foundation models:

- **Foundation Models for Time Series in Conformal Prediction** (Achour et al., 2025) [135]: Achour et al. uncover a structural synergy between zero-shot foundation models and Split Conformal Prediction (SCP). In classical machine learning, limited data must be bifurcated between parameter fitting $\mathcal{D}_{\text{train}}$ and nonconformity calibration $\mathcal{D}_{\text{calib}}$, severely impairing calibration sample size. Because zero-shot foundation models (Lag-Llama, TimesFM, Chronos) require zero target training data, 100% of available in-

domain historical data can be dedicated to calibration. Evaluating across ERCOT, NN5, and M3 benchmarks, Achour et al. show that zero-shot TSFM conformal intervals achieve strictly valid $1 - \alpha$ coverage while producing significantly narrower interval widths than classical GBDT or ARIMA baselines, demonstrating that foundation backbones stabilize post-hoc conformal calibration.

- **Adaptive Conformal Anomaly Detection for Signal Monitoring** (Martinez Gil et al., 2026) [136]: Martinez Gil et al. formulate adaptive conformal anomaly detection for industrial telemetry. Rather than applying arbitrary thresholding heuristics to foundation model reconstruction residuals, the authors utilize weighted quantile conformal prediction bounds to transform residuals into exact, false-alarm-rate controlled p-values. By adaptively updating nonconformity weights online as temporal distributions drift, this framework maintains a guaranteed false alarm rate ($\leq \alpha$) on streaming IoT sensor workloads without retraining foundation backbones.
- **RareCP: Regime-Aware Retrieval for Efficient Conformal Prediction** (Heurich et al., 2026) [137]: Standard conformal inference assumes exchangeability, an assumption

systematically violated by temporal autocorrelation and regime-switching dynamics. Heurich, Granz, and Landgraf introduce RareCP, which decouples smooth distribution drift from distinct macroeconomic/operational regimes. RareCP deploys a Mixture of Experts (MoE) with cosine attention alongside a hypernetwork that retrieves the top- k most relevant historical calibration contexts for any query sequence. By computing asymmetric weighted conformal quantiles over retrieved residuals on the GIFT-Eval benchmark, RareCP preserves nominal coverage guarantees while shrinking interval width by up to 22% compared to Adaptive Conformal Inference (ACI) and raw foundation model quantiles.

Table 10 provides a unified synthesis of the Iteration 9 frontiers, comparing post-training policy optimization architectures, conformal prediction frameworks, physics-informed foundations, and continuous asynchronous event models.

7 Open Challenges and Future Horizons

Despite rapid progress in foundation models and LLM adaptations, temporal machine learning remains in a formative phase. In this section, we delineate six critical open research horizons that must be addressed to achieve truly universal temporal intelligence.

7.1 Irregular Sampling and Continuous-Time Dynamics

Existing foundation models (e.g., Chronos, TimesFM, MOIRAI) predominantly assume regularly sampled, synchronous discrete grids. However, real-world physical telemetry—such as ICU medical records, IoT sensor dropouts, and financial trade ticks—arrives at asynchronous, irregular time intervals. Emerging approaches such as FlowState [41] leverage continuous flow matching and continuous-time neural state-space models to achieve sampling-rate equivariance. Future foundation architectures must natively encode irregular time intervals (Δt) into latent manifolds rather than relying on crude imputation or resampling heuristics.

7.2 Prior-Data Fitted Bayesian Networks

Standard foundation models require thousands of gradient optimization steps during fine-tuning. A compelling alternative is demonstrated by TabPFN-TS [51], which adapts Prior-Data Fitted Networks to time series. By training a transformer exclusively on synthetic prior stochastic processes (e.g., Gaussian processes, Ornstein-Uhlenbeck, and Markov switching models), the network learns to perform Bayesian posterior inference in a single forward pass without parameter updates. Expanding synthetic prior spaces to encompass complex non-linear dynamics represents a fertile research direction.

7.3 Extreme Long-Horizon Context and Serial Scaling

In natural language, modern models handle context windows exceeding 1M tokens. In temporal analysis, historical series often span millions of high-frequency timestamps. Capturing decadal climate shifts, annual retail cycles, and second-level power grid vibrations within a single context window demands extreme context scaling. Architectural innovations such as long-context state-space models (Mamba), hierarchical linear

attention, and serial scaling (as explored in Timer-XL [16] and Timer-S1 [18]) will be essential to eliminate context truncation.

7.4 Unified Multimodal Perception

Real-world decision-making rarely relies on isolated scalar numbers. An industrial operator monitors pressure gauges alongside diagnostic maintenance logs; a macroeconomist reads financial reports while analyzing stock tickers; an environmental scientist tracks satellite imagery across seasonal vegetation cycles [62]. Future foundation models must advance beyond ad-hoc text reprogramming to train native multi-modal architectures from scratch—treating numerical time series, textual reports, spatial optical imagery, and spectral vibrations as co-equal perceptual modalities, as pioneered by Chronicle [98], VLT [97], and ChronoSteer [99].

7.5 Calibrated Uncertainty under Distribution Shift

In safety-critical applications (such as healthcare and grid management), an uncalibrated point forecast is hazardous. As demonstrated by recent reliability audits [26], current zero-shot models frequently underestimate predictive variance when confronted with unprecedented macroeconomic shocks or extreme weather events. Developing conformal prediction techniques, robust Bayesian heads, and flow-based generative sampling that maintain rigorous coverage guarantees under out-of-distribution shifts is imperative.

7.6 Edge Deployment, Hardware Quantization, and Green AI

While research models scale to billions of parameters, industrial IoT environments require deployment on low-power microcontrollers, edge gateways, and automotive Electronic Control Units (ECUs). Systematic cross-platform benchmarking (HoliBench [125]), dynamical systems trajectory sensitivity planning (TQS-PTQ [67]), percentile-based 4-bit activation calibration [68], on-device fully quantized microcontroller training (MCU-FQT [69]), integer mixed-precision quantization (INT8/INT4 on embedded FPGAs [126]), ultra-compact patch encoders [63], [65], and carbon-aware dynamic dispatch (FM-CAC [127]) will be essential to transform temporal foundation research into sustainable, energy-efficient cyber-physical systems.

7.7 Autonomous Living Foundations and Continuous Discovery

With dozens of new preprints posted weekly across machine learning repositories, academic surveys quickly become obsolete. Future research must embrace autonomous, living benchmark harnesses (e.g., LiveHouse-TS [122], Impermanent [130]), continuous automated snowballing pipelines that monitor arXiv and Semantic Scholar streams, vintage-consistent simulation pipelines eliminating lookahead bias (MACROCAST [66]), and automated multi-granularity leakage audits (It’s TIME [117]) to maintain living, contamination-free community standards.

TABLE 10
Taxonomy and Empirical Characterization of Post-Training Policy Optimization (RLHF/RLPF), Distribution-Free Conformal Prediction, Physics Foundations, and Asynchronous Event-Stream Models.

Framework	Key Reference	Core Paradigm	Category &	Policy / Conformal / Physical Mechanism	Target Domain / Hardware	Key Empirical Finding / Theoretical Guarantee
Category A: Reinforcement Learning and Post-Training Policy Optimization (RLHF / RLPF)						
TimeRFT	Li et al. [111] (arXiv’26)	RL Fine-Tuning for TSFMs		Quality-aware temporal rewards + difficulty-aware sample selection	Multi-Domain (TimesFM, Chronos)	Up to -18.6% forecast error under distribution shift; overcomes SFT exposure bias
TimeHF	Qi et al. [112] (arXiv’25)	6B Human-Aligned Foundation		Timeseries Policy Optimization (TPO) with pairwise preference rewards	Enterprise Supply Chain (20k products)	$+33.21\%$ replenishment decision accuracy; aligns asymmetric stockout costs
COUNTS	Parker et al. [113] (NeurIPS’25)	Reasoning LLM for Time Series		RVQ-VAE discrete tokens + Group Relative Policy Optimization (GRPO)	Medical ECG-QA / UCR Archive	Elicits verifiable Chain-of-Thought reasoning; outperforms direct prompting by large margins
TimeMaster	Zhang et al. [114] (arXiv’25)	Multimodal Reasoning via RL		Step-by-step diagnostic rewards over spectrograms and waveforms	Physiological Signals (EMG, ECG, HAR)	Eliminates visual hallucinations; enables structured multimodal diagnostic reasoning
LAST SToP	Gupta et al. [115] (ICML’25)	Asynchronous LLM Adaptation		Language-modeled Asynchronous TS (LASTS) + Stochastic Soft Prompting	Irregular Clinical ICU Telemetry	Outperforms QLoRA on irregular time intervals while preserving pre-trained LLM priors
Category B: Distribution-Free Conformal Prediction and Coverage Guarantees						
TSFM-Conformal	Achour et al. [135] (arXiv’25)	Split Conformal Prediction (SCP)		Zero-shot TSFM base forecasters dedicating 100% target data to calibration	ERCOT, NN5, M3 Benchmarks	Exact $1 - \alpha$ marginal coverage guarantee; produces tighter intervals than GBDTs
Adaptive-CAD	Martinez et al. [136] (ICLR’26)	Conformal Anomaly Detection		Weighted conformal quantile bounds outputting p-values	Industrial IoT Telemetry Streams	Strict false alarm rate control ($\leq \alpha$); stable calibration under non-stationary drift
RareCP	Heurich et al. [137] (arXiv’26)	Regime-Aware Conformal MoE		Cosine-attention retrieval of calibration residuals + hypernetwork drift adaptation	GIFT-Eval Benchmark Suite	Preserves nominal coverage while reducing interval width by up to 22% over ACI
Category C: Physics-Informed Foundations and Continuous Asynchronous Event Streams						
GridSFM	Bhan et al. [82] (arXiv’26)	Physics-Inspired Grid Model		Newton power flow operators with logarithmically penalized slack slacks	Electrical Transmission (10k buses)	2.45% zero-shot generation cost error on unseen grids; $10\times$ solver speedup
Physics-SynHist	Longarini et al. [83] (ICML-W’26)	Cold-Start Pipeline	Physical	Deterministic physical chains (pvlib) generating surrogate histories	Solar PV Facilities (440 sites)	Up to 50% MAE reduction on zero-shot cold start without target fine-tuning
SurF	Rezaei et al. [80] (arXiv’26)	Generative Foundation	Event	Time Rescaling Theorem (TRT) as continuous flow bijection to $\text{Exp}(1)$	Earthquakes, Retweets, commerce	Multi-dataset continuous pre-training; eliminates quadrature integration bottlenecks
TradeFM	Kawawa-Beaudan et al. [81] (arXiv’26)	524M Microstructure	Market	Scale-invariant universal tokenization of limit order books	9,000+ Equities (US, APAC)	Replicates non-linear financial stylized facts; zero-shot international market rollout

7.8 Interactive Multi-Turn Agents and Strategic Tool-Use

Finally, the paradigm of treating time-series intelligence as passive, single-pass point prediction is rapidly yielding to autonomous, tool-augmented agentic ecosystems [100]. Future temporal systems will combine streaming asynchronous perception (TimeInteract [101]), reinforcement-learned sequential decision policies (Cast-R1 [102]), expert trajectory reasoning models (TimeART [103]), domain-specialized constraint verifiers (TS-Reasoner [104]), and automated data cleaning planners (DCATS [105]). Developing formal verification methods for multi-turn agentic actions, tool-use error boundaries, and streaming conversational state retention represents the definitive next frontier for temporal artificial intelligence.

8 Conclusion

In this survey, conducted under the PRISMA 2020 systematic review protocol, we have synthesized the foundational developments at the intersection of Large Language Models and Time Series Foundation Models from 2021 through 2026. We formalized mathematical definitions of temporal sequences, subseries patching, and quantization; introduced a unified five-dimensional taxonomy; and compared 41 core benchmark and foundation models across architecture, tokenization, pre-training corpora, and uncertainty formulations.

Our critical synthesis highlights a decisive architectural demarcation:

- Native TSFMs (such as Chronos, TimesFM, MOIRAI, Timer, and TTM) demonstrate superior parameter efficiency, faster inference latency, and stronger zero-shot nu-

meric forecasting accuracy by training directly on temporal dynamics without linguistic mismatch.

- Repurposed LLMs (such as Time-LLM and multimodal reasoning frameworks) excel not at raw numeric extrapolation, but as cognitive orchestrators—grounding numerical time series within textual narratives, metadata, and decision-support systems.

Looking ahead, the frontier of temporal intelligence lies in resolving benchmark contamination, achieving continuous-time sampling equivariance, calibrating uncertainty under non-stationary shifts, and delivering ultra-efficient edge deployment. We hope this survey provides a rigorous compass for researchers and practitioners building the next generation of universal temporal models.

References

- [1] M. Jin, Y. Kong, Y. Liang, C. Zhang, S. Xue, X. Wang, J. Zhang, Y. Wang, H. Chen, X. Li, V. S. Tseng, Y. Zheng, L. Chen, H. Xiong, S. Pan, and Q. Wen, “Large Models for Time Series and Spatio-Temporal Data: A Survey and Outlook,” *ACM Computing Surveys* 2024, 2023. [Online]. Available: <https://arxiv.org/abs/2310.10196>
- [2] Y. Nie, N. H. Nguyen, P. Sinthong, and J. Kalagnanam, “A Time Series is Worth 64 Words: Long-term Forecasting with Transformers,” in *ICLR* 2023, 2022. [Online]. Available: <https://arxiv.org/abs/2211.14730>
- [3] A. F. Ansari, L. Stella, C. Turkmen, X. Zhang, P. Mercado, H. Shen, O. Shchur, S. S. Rangapuram, S. P. Arango, S. Kapoor, J. Zschneider, D. C. Maddix, H. Wang, M. W. Mahoney, K. Torkkola, A. G. Wilson, M. Bohlke-Schneider, and Y. Wang, “Chronos: Learning the Language of Time Series,” in *ICML 2024 / TMLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2403.07815>
- [4] A. Das, W. Kong, R. Sen, and Y. Zhou, “A decoder-only foundation model for time-series forecasting,” in *ICML* 2024, 2023. [Online]. Available: <https://arxiv.org/abs/2310.10688>
- [5] G. Woo, C. Liu, A. Kumar, C. Xiong, S. Savarese, and D. Sahoo, “Unified Training of Universal Time Series Forecasting Transformers,” in *ICML* 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2402.02592>
- [6] N. Gruver, M. Finzi, S. Qiu, and A. G. Wilson, “Large Language Models Are Zero-Shot Time Series Forecasters,” in *NeurIPS* 2023, 2023. [Online]. Available: <https://arxiv.org/abs/2310.07820>
- [7] M. Tan, M. A. Merrill, V. Gupta, T. Althoff, and T. Hartvigsen, “Are Language Models Actually Useful for Time Series Forecasting?” in *NeurIPS* 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2406.16964>
- [8] A. F. Ansari, O. Shchur, J. Küken, A. Auer, B. Han, P. Mercado, S. S. Rangapuram, H. Shen, L. Stella, X. Zhang, M. Goswami, S. Kapoor, D. C. Maddix, P. Gueron, T. Hu, J. Yin, N. Erickson, P. M. Desai, H. Wang, H. Rangwala, G. Karypis, Y. Wang, and M. Bohlke-Schneider, “Chronos-2: From Univariate to Universal Forecasting,” *arXiv preprint arXiv:2510.15821*, 2025. [Online]. Available: <https://arxiv.org/abs/2510.15821>
- [9] A. Das, M. Faw, R. Sen, and Y. Zhou, “In-Context Fine-Tuning for Time-Series Foundation Models,” in *NeurIPS* 2024 Workshop, 2024. [Online]. Available: <https://arxiv.org/abs/2410.24087>
- [10] X. Liu, J. Liu, G. Woo, T. Aksu, Y. Liang, R. Zimmermann, C. Liu, S. Savarese, C. Xiong, and D. Sahoo, “Moirai-MoE: Empowering Time Series Foundation Models with Sparse Mixture of Experts,” *arXiv preprint arXiv:2410.10469*, 2024. [Online]. Available: <https://arxiv.org/abs/2410.10469>
- [11] C. Liu, T. Aksu, J. Liu, X. Liu, H. Yan, Q. Pham, S. Savarese, D. Sahoo, C. Xiong, and J. Li, “Moirai 2.0: When Less Is More for Time Series Forecasting,” *arXiv preprint arXiv:2511.11698*, 2025. [Online]. Available: <https://arxiv.org/abs/2511.11698>
- [12] M. Goswami, K. Szafer, A. Choudhry, Y. Cai, S. Li, and A. Dubrawski, “MOMENT: A Family of Open Time-series Foundation Models,” in *ICML* 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2402.03885>
- [13] V. Ekambaram, A. Jati, P. Dayama, S. Mukherjee, N. H. Nguyen, W. M. Gifford, C. Reddy, and J. Kalagnanam, “Tiny Time Mixers (TTMs): Fast Pre-trained Models for Enhanced Zero/Few-Shot Forecasting of Multivariate Time Series,” in *NeurIPS* 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2401.03955>
- [14] K. Rasul, A. Ashok, A. R. Williams, H. Ghonia, R. Bhagwatkar, A. Khorasani, M. J. D. Bayazi, G. Adamopoulos, R. Riachi, N. Hassen, M. Biloš, S. Garg, A. Schneider, N. Chapados, A. Drouin, V. Zantedeschi, Y. Nevmyvaka, and I. Rish, “Lag-Llama: Towards Foundation Models for Probabilistic Time Series Forecasting,” in *ICML 2024 Workshop*, 2023. [Online]. Available: <https://arxiv.org/abs/2310.08278>
- [15] Y. Liu, H. Zhang, C. Li, X. Huang, J. Wang, and M. Long, “Timer: Generative Pre-trained Transformers Are Large Time Series Models,” in *ICML* 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2402.02368>
- [16] Y. Liu, G. Qin, X. Huang, J. Wang, and M. Long, “Timer-XL: Long-Context Transformers for Unified Time Series Forecasting,” in *NeurIPS* 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2410.04803>
- [17] Y. Liu, G. Qin, Z. Shi, Z. Chen, C. Yang, X. Huang, J. Wang, and M. Long, “Sundial: A Family of Highly Capable Time Series Foundation Models,” *arXiv preprint arXiv:2502.00816*, 2025. [Online]. Available: <https://arxiv.org/abs/2502.00816>
- [18] Y. Liu, X. Su, S. Wang, H. Zhang, H. Liu, Y. Wang, Z. Ye, Y. Xiang, J. Wang, and M. Long, “Timer-S1: A Billion-Scale Time Series Foundation Model with Serial Scaling,” *arXiv preprint arXiv:2603.04791*, 2026. [Online]. Available: <https://arxiv.org/abs/2603.04791>
- [19] M. Jin, S. Wang, L. Ma, Z. Chu, J. Y. Zhang, X. Shi, P.-Y. Chen, Y. Liang, Y.-F. Li, S. Pan, and Q. Wen, “Time-LLM: Time Series Forecasting by Reprogramming Large Language Models,” in *ICLR* 2024, 2023. [Online]. Available: <https://arxiv.org/abs/2310.01728>
- [20] T. Zhou, P. Niu, X. Wang, L. Sun, and R. Jin, “One Fits All: Power General Time Series Analysis by Pretrained LM,” in *NeurIPS* 2023, 2023. [Online]. Available: <https://arxiv.org/abs/2302.11939>
- [21] D. Cao, F. Jia, S. O. Arik, T. Pfister, Y. Zheng, W. Ye, and Y. Liu, “TEMPO: Prompt-based Generative Pre-trained Transformer for Time Series Forecasting,” in *ICLR* 2024, 2023. [Online]. Available: <https://arxiv.org/abs/2310.04948>
- [22] Y. Kong, Y. Yang, Y. Hwang, W. Du, S. Zohren, Z. Wang, M. Jin, and Q. Wen, “Time-MQA: Time Series Multi-Task Question Answering with Context Enhancement,” in *ACL* 2025, 2025. [Online]. Available: <https://arxiv.org/abs/2503.01875>
- [23] T. Guan, Z. Meng, D. Li, S. Wang, C.-H. H. Yang, Q. Wen, Z. Liu, S. M. Siniscalchi, M. Jin, and S. Pan, “TimeOmni-1: Incentivizing Complex Reasoning with Time Series in Large Language Models,” in *ICLR* 2026, 2025. [Online]. Available: <https://arxiv.org/abs/2509.24803>
- [24] T. Aksu, G. Woo, J. Liu, X. Liu, C. Liu, S. Savarese, C. Xiong, and D. Sahoo, “GIFT-Eval: A Benchmark For General Time Series Forecasting Model Evaluation,” in *NeurIPS* 2024 D&B, 2024. [Online]. Available: <https://arxiv.org/abs/2410.10393>
- [25] M. Meyer, S. Kaltenpoth, K. Zalipski, and O. Müller, “Rethinking Evaluation in the Era of Time Series Foundation Models: (Un)known Information Leakage Challenges,” *arXiv preprint arXiv:2510.13654*, 2025. [Online]. Available: <https://arxiv.org/abs/2510.13654>
- [26] P. Michael, M. Symeonides, and D. Trihinas, “Evaluating Accuracy and Probabilistic Reliability of Zero-Shot Time Series Foundation Models,” *arXiv preprint arXiv:2609.25788*, 2026. [Online]. Available: <https://arxiv.org/abs/2609.25788>
- [27] S. Dooley, G. S. Khurana, C. Mohapatra, S. Naidu, and C. White, “ForecastPFN: Synthetically-Trained Zero-Shot Forecasting,” in *NeurIPS* 2023, 2023. [Online]. Available: <https://arxiv.org/abs/2311.01933>
- [28] A. Auer, P. Podest, D. Klotz, S. Böck, G. Klambauer, and S. Hochreiter, “TiRex: Zero-Shot Forecasting Across Long and Short Horizons with Enhanced In-Context Learning,” *arXiv preprint arXiv:2505.23719*, 2025. [Online]. Available: <https://arxiv.org/abs/2505.23719>
- [29] X. Wu, H. Cheng, X. Qiu, Z. Li, J. Hu, C. Guo, and B. Yang, “FLAME: Flow Enhanced Legendre Memory

- Models for General Time Series Forecasting,” arXiv preprint arXiv:2512.14253, 2025. [Online]. Available: <https://arxiv.org/abs/2512.14253>
- [30] L. Meyer, C. Sole, H. Xiang, N. Li, L. Franceschino, A. Quera-Bofarull, M. P. Scholl, J. Fainberg, and G. Négier, “ t_0 : A Time-Series Foundation Model for Forecasting with Context,” arXiv preprint arXiv:2609.24559, 2026. [Online]. Available: <https://arxiv.org/abs/2609.24559>
- [31] C. Sun, H. Li, Y. Li, and S. Hong, “TEST: Text Prototype Aligned Embedding to Activate LLM’s Ability for Time Series,” in NeurIPS 2024, 2023. [Online]. Available: <https://arxiv.org/abs/2308.08241>
- [32] P. Liu, H. Guo, T. Dai, N. Li, J. Bao, X. Ren, Y. Jiang, and S.-T. Xia, “CALF: Aligning LLMs for Time Series Forecasting via Cross-modal Fine-Tuning,” in KDD 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2403.07300>
- [33] X. Liu, J. Hu, Y. Li, S. Diao, Y. Liang, B. Hooi, and R. Zimmermann, “UniTime: A Language-Empowered Unified Model for Cross-Domain Time Series Forecasting,” arXiv preprint arXiv:2310.09751, 2023. [Online]. Available: <https://arxiv.org/abs/2310.09751>
- [34] M. Chen, L. Shen, Z. Li, X. J. Wang, J. Sun, and C. Liu, “VisionTS: Visual Masked Autoencoders Are Free-Lunch Zero-Shot Time Series Forecasters,” in NeurIPS 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2408.17253>
- [35] S. Zhong, W. Ruan, M. Jin, H. Li, Q. Wen, and Y. Liang, “Time-VLM: Exploring Multimodal Vision-Language Models for Augmented Time Series Forecasting,” in ICML 2025, 2025. [Online]. Available: <https://arxiv.org/abs/2502.04395>
- [36] W. Wu, Z. Zhang, L. Liu, X. Xu, J. Zhuang, K. Fan, Q. Lv, J. Liu, C. Zhang, Z. Yuan, S. Hou, T. Lin, K. Chen, B. Zhou, and C. Zhang, “SciTS: Scientific Time Series Understanding and Generation with LLMs,” arXiv preprint arXiv:2510.03255, 2025. [Online]. Available: <https://arxiv.org/abs/2510.03255>
- [37] Y. Zhang, Y. Zhang, M. Zheng, K. Chen, C. Gao, R. Ge, S. Teng, A. Jelloul, J. Rao, X. Guo, C.-W. Fang, Z. Zheng, and J. Yang, “Insight Miner: A Time Series Analysis Dataset for Cross-Domain Alignment with Natural Language,” arXiv preprint arXiv:2512.11251, 2025. [Online]. Available: <https://arxiv.org/abs/2512.11251>
- [38] H. Liu, S. Xu, Z. Zhao, L. Kong, H. Kamarthi, A. B. Sasanur, M. Sharma, J. Cui, Q. Wen, C. Zhang, and B. A. Prakash, “Time-MMD: Multi-Domain Multimodal Dataset for Time Series Analysis,” in NeurIPS 2024 D&B, 2024. [Online]. Available: <https://arxiv.org/abs/2406.08627>
- [39] S. Gao, T. Koker, O. Queen, T. Hartvigsen, T. Tsiligkaridis, and M. Zitnik, “UniTS: A Unified Multi-Task Time Series Model,” in NeurIPS 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2403.00131>
- [40] X. Shi, S. Wang, Y. Nie, D. Li, Z. Ye, Q. Wen, and M. Jin, “Time-MoE: Billion-Scale Time Series Foundation Models with Mixture of Experts,” in ICLR 2025, 2024. [Online]. Available: <https://arxiv.org/abs/2409.16040>
- [41] L. Graf, T. Ortner, S. Woźniak, and A. Pantazi, “FlowState: Sampling-Rate-Equivariant Time-Series Forecasting,” arXiv preprint arXiv:2508.05287, 2025. [Online]. Available: <https://arxiv.org/abs/2508.05287>
- [42] H. Xue and F. D. Salim, “PromptCast: A New Prompt-based Learning Paradigm for Time Series Forecasting,” IEEE TKDE 2023, 2022. [Online]. Available: <https://arxiv.org/abs/2210.08964>
- [43] S. Wang, H. Wu, X. Shi, T. Hu, H. Luo, L. Ma, J. Y. Zhang, and J. Zhou, “TimeMixer: Decomposable Multiscale Mixing for Time Series Forecasting,” in ICLR 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2405.14616>
- [44] S. Wang, J. Li, X. Shi, Z. Ye, B. Mo, W. Lin, S. Ju, Z. Chu, and M. Jin, “TimeMixer++: A General Time Series Pattern Machine for Universal Predictive Analysis,” in ICLR 2025, 2024. [Online]. Available: <https://arxiv.org/abs/2410.16032>
- [45] E. Khwaja, C. Lettieri, G. Woo, E. Belouadah, M. Cenac, G. Jarry, E. Paquin, X. Zhao, V. Zhukov, O. Abou-Amal, C. Liu, A. Talwalkar, and D. Asker, “Toto 2.0: Time Series Forecasting Enters the Scaling Era,” arXiv preprint arXiv:2605.20119, 2026. [Online]. Available: <https://arxiv.org/abs/2605.20119>
- [46] Y. Liu, G. Qin, X. Huang, J. Wang, and M. Long, “AutoTimes: Autoregressive Time Series Forecasters via Large Language Models,” in NeurIPS 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2402.02370>
- [47] R. Tacconelli, “Cadence: Error-Bounded Lossy Compression of Demand Time Series with a Time-Series Foundation Model,” arXiv preprint arXiv:2609.06008, 2026. [Online]. Available: <https://arxiv.org/abs/2609.06008>
- [48] X. Wang, T. Zhou, J. Gao, B. Ding, and J. Zhou, “Output Scaling: YingLong-Delayed Chain of Thought in a Large Pretrained Time Series Forecasting Model,” arXiv preprint arXiv:2506.11029, 2025. [Online]. Available: <https://arxiv.org/abs/2506.11029>
- [49] B. Cohen, E. Khwaja, Y. Doubli, S. Lemaachi, C. Lettieri, C. Masson, H. Miccinilli, E. Ramé, Q. Ren, A. Rostamizadeh, J. O. du Terrail, A.-M. Toon, K. Wang, S. Xie, Z. Xu, V. Zhukova, D. Asker, A. Talwalkar, and O. Abou-Amal, “This Time is Different: An Observability Perspective on Time Series Foundation Models,” arXiv preprint arXiv:2505.14766, 2025. [Online]. Available: <https://arxiv.org/abs/2505.14766>
- [50] S. Xie, B. Abdessalem, Z. Xiao, Y. A. E. Hili, A. Odonnat, Z. Dong, L. Zan, T. Palpanas, J. Zhang, L. Pan, K. Zhang, and M. Tiomoko, “Tabby: An Open Pretraining Recipe for Time Series Foundation Models,” arXiv preprint arXiv:2609.13956, 2026. [Online]. Available: <https://arxiv.org/abs/2609.13956>
- [51] S. B. Hoo, S. Müller, D. Salinas, and F. Hutter, “From Tables to Time: Extending TabPFN-v2 to Time Series Forecasting,” arXiv preprint arXiv:2501.02945, 2025. [Online]. Available: <https://arxiv.org/abs/2501.02945>
- [52] J. Lu, Y. Sun, and S. Yang, “In-context Time Series Predictor,” in ICLR 2025, 2024. [Online]. Available: <https://arxiv.org/abs/2405.14982>
- [53] Y. Wang, H. Wu, J. Dong, G. Qin, H. Zhang, Y. Liu, Y. Qiu, J. Wang, and M. Long, “TimeXer: Empowering Transformers for Time Series Forecasting with Exogenous Variables,” in NeurIPS 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2402.19072>
- [54] K. Feng, S. Lan, Y. Fang, W. He, S. Lu, S. Gu, L. Ma, X. Lu, and K. Ren, “Kairos: Toward Adaptive and Parameter-Efficient Time Series Foundation Models,” arXiv preprint arXiv:2509.25826, 2025. [Online]. Available: <https://arxiv.org/abs/2509.25826>
- [55] Y. Wang, Y. Qiu, P. Chen, Y. Shu, Z. Rao, L. Pan, B. Yang, and C. Guo, “LightGTS: A Lightweight General Time Series Forecasting Model,” in ICML 2025, 2025. [Online]. Available: <https://arxiv.org/abs/2506.06005>
- [56] X. Zhou, Q. Yao, Y. Zhao, C. Liu, F. Salim, X. Yuan, Y. Wen, and M. Jin, “EIDOS: Latent-Space Predictive Learning for Time Series Foundation Models,” arXiv preprint arXiv:2602.14024, 2026. [Online]. Available: <https://arxiv.org/abs/2602.14024>
- [57] A. Goodge, W. S. Ng, B. Hooi, and S. K. Ng, “Spatio-Temporal Foundation Models: Vision, Challenges, and Opportunities,” arXiv preprint arXiv:2501.09045, 2025. [Online]. Available: <https://arxiv.org/abs/2501.09045>
- [58] Z. Li, L. Xia, L. Shi, Y. Xu, D. Yin, and C. Huang, “OpenCity: Open Spatio-Temporal Foundation Models for Traffic Prediction,” in NeurIPS 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2408.10269>
- [59] Y. Yuan, J. Ding, J. Feng, D. Jin, and Y. Li, “UniST: A Prompt-Empowered Universal Model for Urban Spatio-Temporal Prediction,” in KDD 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2402.11838>
- [60] Y. Yuan, C. Han, J. Ding, G. Zhang, D. Jin, and Y. Li, “Diffusion Transformers as Open-World Spatiotemporal Foundation Models,” in NeurIPS 2025, 2024. [Online]. Available: <https://arxiv.org/abs/2411.12164>
- [61] W. Chen, Y. Wu, J. Chen, X. Zhou, and Y. Liang, “UrbanFM: Scaling Urban Spatio-Temporal Foundation Models,” arXiv preprint arXiv:2602.20677, 2026. [Online]. Available: <https://arxiv.org/abs/2602.20677>
- [62] X. Qin, D. Wang, J. Zhang, F. Wang, X. Su, B. Du, and L. Zhang, “TiMo: Spatiotemporal Foundation Model for Satellite Image Time Series,” arXiv preprint arXiv:2505.08723, 2025. [Online]. Available: <https://arxiv.org/abs/2505.08723>
- [63] F. Birkel, “Tiny-TSM: Efficiently Training a Lightweight SOTA Time Series Foundation Model,” arXiv preprint arXiv:2511.19272, 2025. [Online]. Available: <https://arxiv.org/abs/2511.19272>

- [64] S. Pradhan, N. Bahadori, and P. Amini, "APEX: A Network-Native Time-Series Foundation Model for Forecasting and Anomaly Detection for Wireless Edge Operations," arXiv preprint arXiv:2606.11553, 2026. [Online]. Available: <https://arxiv.org/abs/2606.11553>
- [65] M. Cheraghinia, E. D. Poorter, J. Fontaine, K. S. Kim, M. Debbah, and A. Shahid, "Lightweight Foundation Model for Wireless Time Series Downstream Tasks on Edge Devices," arXiv preprint arXiv:2511.14895, 2025. [Online]. Available: <https://arxiv.org/abs/2511.14895>
- [66] A. Carriero, D. Pettenuzzo, and S. Shekhar, "MACROCAST: A Vintage-Consistent Time Series Foundation Model for Real-Time Macroeconomic Forecasting," arXiv preprint arXiv:2606.28670, 2026. [Online]. Available: <https://arxiv.org/abs/2606.28670>
- [67] M. Pavlova, H. B. H. Zhu, L. Vitanova, E. Semenova, and Y. Li, "Quantizing Time-Series Models As Dynamical Systems: Trajectory-Based Quantization Sensitivity Score," arXiv preprint arXiv:2606.13300, 2026. [Online]. Available: <https://arxiv.org/abs/2606.13300>
- [68] J. Ye and I. G. Wanjiku, "Calibration Bets on the Past: Post-Training Quantization for Financial Time-Series Forecasting," arXiv preprint arXiv:2608.12259, 2026. [Online]. Available: <https://arxiv.org/abs/2608.12259>
- [69] M. Deutel, F. Hannig, C. Mutschler, and J. Teich, "On-Device Training of Fully Quantized Deep Neural Networks on Cortex-M Microcontrollers," arXiv preprint arXiv:2407.10734, 2024. [Online]. Available: <https://arxiv.org/abs/2407.10734>
- [70] M. Gorbett, H. Shirazi, and I. Ray, "Sparse Binary Transformers for Multivariate Time Series Modeling," arXiv preprint arXiv:2308.04637, 2023. [Online]. Available: <https://arxiv.org/abs/2308.04637>
- [71] R. Yang, H. Zhu, Q. Gao, Y. Ma, H. Wei, and K. Wen, "Q-DEQ: Discrete Solving and Quantization for Deep Equilibrium Models in Time Series Forecasting under Edge Deployment Coding Constraints," arXiv preprint arXiv:2609.24042, 2026. [Online]. Available: <https://arxiv.org/abs/2609.24042>
- [72] X.-Y. Hu, S. Liang, C. Feng, and S.-Q. Zhang, "SGA: Uncertainty Quantification for Multi-Step Forecasting in Time Series Foundation Models," arXiv preprint arXiv:2609.28582, 2026. [Online]. Available: <https://arxiv.org/abs/2609.28582>
- [73] H. Kim, S. Kim, J. Mok, and S. Yoon, "Battling the Non-stationarity in Time Series Forecasting via Test-time Adaptation," arXiv preprint arXiv:2501.04970, 2025. [Online]. Available: <https://arxiv.org/abs/2501.04970>
- [74] T. Dang, S. Chatterjee, H. Jia, Y. Wu, F. Salim, and F. Kawsar, "AdaNODEs: Test Time Adaptation for Time Series Forecasting Using Neural ODEs," arXiv preprint arXiv:2601.12893, 2026. [Online]. Available: <https://arxiv.org/abs/2601.12893>
- [75] I. Kumar, A. Tiwari, S. K. Jasti, and A. H. Lade, "RG-TTA: Regime-Guided Meta-Control for Test-Time Adaptation in Streaming Time Series," arXiv preprint arXiv:2603.27814, 2026. [Online]. Available: <https://arxiv.org/abs/2603.27814>
- [76] M. Jander, W. van Heeswijk, and M. Mes, "Causal Analysis for Time Series Foundation Models," arXiv preprint arXiv:2608.24303, 2026. [Online]. Available: <https://arxiv.org/abs/2608.24303>
- [77] G. Stein, M. Shadaydeh, and J. Denzler, "Embracing the black box: Heading towards foundation models for causal discovery from time series data," arXiv preprint arXiv:2402.09305, 2024. [Online]. Available: <https://arxiv.org/abs/2402.09305>
- [78] D. Thumm and Y. Chen, "Interventional Time Series Priors for Causal Foundation Models," arXiv preprint arXiv:2603.11090, 2026. [Online]. Available: <https://arxiv.org/abs/2603.11090>
- [79] Y. Xia, C. Xu, Y. Liang, L. Zhao, Q. Wen, R. Zimmermann, and J. Bian, "Causal Time Series Generation via Diffusion Models," arXiv preprint arXiv:2509.20846, 2025. [Online]. Available: <https://arxiv.org/abs/2509.20846>
- [80] M. R. Rezaei, T. Balaji, and R. G. Krishnan, "SurF: A Generative Model for Multivariate Irregular Time Series Forecasting," arXiv preprint arXiv:2605.14069, 2026. [Online]. Available: <https://arxiv.org/abs/2605.14069>
- [81] M. Kawawa-Beaudan, S. Sood, K. Papasotiriou, D. Borrajo, and M. Veloso, "TradeFM: A Generative Foundation Model for Trade-flow and Market Microstructure," arXiv preprint arXiv:2602.23784, 2026. [Online]. Available: <https://arxiv.org/abs/2602.23784>
- [82] L. Bhan, W. Yang, M. Capetz, and B. Zhang, "GridSFM: A Foundation Model for Solving AC Optimal Power Flow," arXiv preprint arXiv:2609.30173, 2026. [Online]. Available: <https://arxiv.org/abs/2609.30173>
- [83] L. Longarini, A. Rongoni, S. Silenzi, E. Frontoni, and R. Rosati, "Time series Foundation Models based on Physics-Informed Synthetic Histories for Cold-Start Photovoltaic Forecasting," in ICML 2026 Workshop, 2026. [Online]. Available: <https://arxiv.org/abs/2606.07457>
- [84] C. Chang, W.-Y. Wang, W.-C. Peng, and T.-F. Chen, "LLM4TS: Aligning Pre-Trained LLMs as Data-Efficient Time-Series Forecasters," arXiv preprint arXiv:2308.08469, 2023. [Online]. Available: <https://arxiv.org/abs/2308.08469>
- [85] C. Liu, Q. Xu, H. Miao, S. Yang, L. Zhang, C. Long, Z. Li, and R. Zhao, "TimeCMA: Towards LLM-Empowered Multivariate Time Series Forecasting via Cross-Modality Alignment," arXiv preprint arXiv:2406.01638, 2024. [Online]. Available: <https://arxiv.org/abs/2406.01638>
- [86] Z. Pan, Y. Jiang, S. Garg, A. Schneider, Y. Nevmyvaka, and D. Song, "S²IP-LLM: Semantic Space Informed Prompt Learning with LLM for Time Series Forecasting," arXiv preprint arXiv:2403.05798, 2024. [Online]. Available: <https://arxiv.org/abs/2403.05798>
- [87] M. Kowsher, M. S. I. Sobuj, N. J. Prottasha, E. A. Alanis, O. O. Garibay, and N. Yousefi, "LLM-Mixer: Multiscale Mixing in LLMs for Time Series Forecasting," arXiv preprint arXiv:2410.11674, 2024. [Online]. Available: <https://arxiv.org/abs/2410.11674>
- [88] J. Liang, X. Gu, X. Xie, Y. Wu, and X. Wang, "TAC-Time: Texts as Channels For Multimodal Time Series Forecasting," arXiv preprint arXiv:2609.24156, 2026. [Online]. Available: <https://arxiv.org/abs/2609.24156>
- [89] Z. Xie, Z. Li, X. He, L. Xu, X. Wen, T. Zhang, J. Chen, R. Shi, and D. Pei, "ChatTS: Aligning Time Series with LLMs via Synthetic Data for Enhanced Understanding and Reasoning," arXiv preprint arXiv:2412.03104, 2024. [Online]. Available: <https://arxiv.org/abs/2412.03104>
- [90] C. Wang, Q. Qi, J. Wang, H. Sun, Z. Zhuang, J. Wu, L. Zhang, and J. Liao, "ChatTime: A Unified Multimodal Time Series Foundation Model Bridging Numerical and Textual Data," arXiv preprint arXiv:2412.11376, 2024. [Online]. Available: <https://arxiv.org/abs/2412.11376>
- [91] T. Guan, S. Pan, J. Barthelemy, Z. Li, Y. Cai, C. Alippi, M. Jin, and S. Pan, "TimeOmni-VL: Unified Models for Time Series Understanding and Generation," in ICML 2026, 2026. [Online]. Available: <https://arxiv.org/abs/2602.17149>
- [92] Z. Kan, Y. Chen, K. Li, A. Wen, X. Wang, L. Wang, J. Duan, S. Wang, H. Liu, and T. Chen, "TRACE: A Temporal Conditional Estimation for Multimodal Time Series Foundation Models," arXiv preprint arXiv:2606.06285, 2026. [Online]. Available: <https://arxiv.org/abs/2606.06285>
- [93] P. Langer, T. Kaar, M. Rosenblatt, M. A. Xu, W. Chow, M. Maritsch, R. Jakob, N. Wang, J. Liu, A. Verma, B. Han, D. S. Kim, H. Chubb, S. Ceresnak, A. Zahedivash, A. T. S. Sandhu, F. Rodriguez, D. McDuff, E. Fleisch, O. Aalami, F. Barata, and P. Schmiedmayer, "OpenTSLM: Time-Series Language Models for Reasoning over Multivariate Medical Text- and Time-Series Data," arXiv preprint arXiv:2510.02410, 2025. [Online]. Available: <https://arxiv.org/abs/2510.02410>
- [94] Z. Li, L. Xia, J. Tang, Y. Xu, L. Shi, L. Xia, D. Yin, and C. Huang, "UrbanGPT: Spatio-Temporal Large Language Models," in KDD 2024, 2024. [Online]. Available: <https://arxiv.org/abs/2403.00813>
- [95] C. Liu, S. Yang, Q. Xu, Z. Li, C. Long, Z. Li, and R. Zhao, "Spatial-Temporal Large Language Model for Traffic Prediction," TKDE 2025, 2024. [Online]. Available: <https://arxiv.org/abs/2401.10134>
- [96] L. Shen, M. Chen, X. Liu, H. Fu, X. Ren, J. Sun, Z. Li, and C. Liu, "VisionTS++: Cross-Modal Time Series Foundation Model with Continual Pre-trained Vision Backbones," arXiv preprint arXiv:2508.04379, 2025. [Online]. Available: <https://arxiv.org/abs/2508.04379>
- [97] H. Wang, J. Yan, X. Wang, and L. Ren, "VLT: A Vision-Language-Time Series Multimodal Foundation Model

- for Industrial Intelligence,” arXiv preprint arXiv:2607.14510, 2026. [Online]. Available: <https://arxiv.org/abs/2607.14510>
- [98] P. Quinlan, J. Levasseur, Q. Li, and X. Zhu, “Chronicle: A Multimodal Foundation Model for Joint Language and Time Series Understanding,” arXiv preprint arXiv:2605.20268, 2026. [Online]. Available: <https://arxiv.org/abs/2605.20268>
- [99] C. Wang, Q. Qi, Z. Rao, L. Pan, and J. Wang, “ChronoSteer: Bridging Large Language Model and Time Series Foundation Model via Synthetic Cross-Modal Alignment Dataset,” arXiv preprint arXiv:2505.10083, 2025. [Online]. Available: <https://arxiv.org/abs/2505.10083>
- [100] Y. Kong, Q. Yao, Y. Nie, Y. Li, Y. Shao, S. Zohren, A. Vettoruzzo, J. Vanschoren, M. Jin, and Q. Wen, “TimeSageMT: A Multi-Turn Benchmark for Evaluating Agentic Time Series Reasoning,” arXiv preprint arXiv:2606.01498, 2026. [Online]. Available: <https://arxiv.org/abs/2606.01498>
- [101] S. Pan, Y. Gu, Y. Guo, W. Jin, B. Du, S. Pan, and M. Jin, “TimeInteract: Towards Real-Time Interactive Intelligence for Streaming Time Series,” arXiv preprint arXiv:2609.26389, 2026. [Online]. Available: <https://arxiv.org/abs/2609.26389>
- [102] X. Tao, M. Cheng, C. Jiang, T. Gao, H. Zhang, and Y. Liu, “Cast-R1: Learning Tool-Augmented Sequential Decision Policies for Time Series Forecasting,” arXiv preprint arXiv:2602.13802, 2026. [Online]. Available: <https://arxiv.org/abs/2602.13802>
- [103] X. Wu, J. Lu, Z. Li, X. Qiu, J. Hu, C. Guo, C. S. Jensen, and B. Yang, “TimeART: Towards Agentic Time Series Reasoning via Tool-Augmentation,” arXiv preprint arXiv:2601.13653, 2026. [Online]. Available: <https://arxiv.org/abs/2601.13653>
- [104] W. Ye, W. Yang, D. Cao, Y. Zhang, L. Tang, J. Cai, and Y. Liu, “TS-Reasoner: Domain-Oriented Time Series Inference Agents for Reasoning and Automated Analysis,” arXiv preprint arXiv:2410.04047, 2024. [Online]. Available: <https://arxiv.org/abs/2410.04047>
- [105] C.-C. M. Yeh, V. Lai, U. S. Saini, X. Fan, Y. Fan, J. Wang, X. Dai, and Y. Zheng, “Empowering Time Series Forecasting with LLM-Agents,” arXiv preprint arXiv:2508.04231, 2025. [Online]. Available: <https://arxiv.org/abs/2508.04231>
- [106] J. Kang, S. Han, H. Choe, and S. Lee, “Traceable Multi-Agent System for Knowledge-Based Forecasting,” arXiv preprint arXiv:2608.03339, 2026. [Online]. Available: <https://arxiv.org/abs/2608.03339>
- [107] P. Trirat, J. M. Kwak, J. Heo, H. Lee, and S. J. Hwang, “Multimodal Collaborative Debate for Zero-Shot Time Series Reasoning,” arXiv preprint arXiv:2601.19151, 2026. [Online]. Available: <https://arxiv.org/abs/2601.19151>
- [108] J. Yang, Y. Zheng, J. Sun, X. Fan, J. Wang, L. Wang, Z. Xu, Q. Liu, Z. Fang, Y. Cai, and P. S. Yu, “TimeEvo: Failure-Driven Self-Evolution of a Time Series Agent,” arXiv preprint arXiv:2609.27277, 2026. [Online]. Available: <https://arxiv.org/abs/2609.27277>
- [109] Y. Hu, X. Dai, Z. Qu, Y. Liu, Z. Dong, J. ming Yang, and Q. Xu, “When Tomorrow Becomes Today: Self-Evolving Policies for Agentic Time-Series Forecasting,” arXiv preprint arXiv:2609.24862, 2026. [Online]. Available: <https://arxiv.org/abs/2609.24862>
- [110] C. Shen, W. Yu, F. Wu, D. Song, H. Tong, D. Luo, W. Cheng, H. Chen, and J. Ni, “MetaCaster: Meta-Harness-Optimized Agent for End-to-End Few-Shot Learning of Lightweight Time Series Forecasters,” arXiv preprint arXiv:2608.23473, 2026. [Online]. Available: <https://arxiv.org/abs/2608.23473>
- [111] S. Li, Y. Chen, Z. Zhu, Y. Pan, Y. Guo, M. Huang, and H. Xiong, “TimeRFT: Stimulating Generalizable Time Series Forecasting for TSFMs via Reinforcement Finetuning,” arXiv preprint arXiv:2605.00015, 2026. [Online]. Available: <https://arxiv.org/abs/2605.00015>
- [112] Y. Qi, H. Hu, D. Lei, J. Zhang, Z. Shi, Y. Huang, Z. Chen, X. Lin, and Z.-J. M. Shen, “TimeHF: Billion-Scale Time Series Models Guided by Human Feedback,” arXiv preprint arXiv:2501.15942, 2025. [Online]. Available: <https://arxiv.org/abs/2501.15942>
- [113] F. Parker, N. Chan, C. Zhang, and K. Ghobadi, “Eliciting Chain-of-Thought Reasoning for Time Series Analysis using Reinforcement Learning,” in NeurIPS 2025, 2025. [Online]. Available: <https://arxiv.org/abs/2510.01116>
- [114] J. Zhang, L. Feng, X. Guo, Y. Wu, Y. Dong, and D. Xu, “TimeMaster: Training Time-Series Multimodal LLMs to Reason via Reinforcement Learning,” arXiv preprint arXiv:2506.13705, 2025. [Online]. Available: <https://arxiv.org/abs/2506.13705>
- [115] S. Gupta, T. Durand, G. Taylor, and L. W. Bialokozowicz, “LAST SToP For Modeling Asynchronous Time Series,” in ICML 2025, 2025. [Online]. Available: <https://arxiv.org/abs/2502.01922>
- [116] O. Shchur, A. F. Ansari, C. Turkmen, L. Stella, N. Erickson, P. Gueron, M. Bohlke-Schneider, and Y. Wang, “fev-bench: A Realistic Benchmark for Time Series Forecasting,” arXiv preprint arXiv:2509.26468, 2025. [Online]. Available: <https://arxiv.org/abs/2509.26468>
- [117] Z. Qiao, S. Pan, A. Wang, V. Zhukova, Y. Liu, X. Jiang, Q. Wen, M. Long, M. Jin, and C. Liu, “It’s TIME: Towards the Next Generation of Time Series Forecasting Benchmarks,” arXiv preprint arXiv:2602.12147, 2026. [Online]. Available: <https://arxiv.org/abs/2602.12147>
- [118] P. Chen, Z. Zhuang, H. Chen, J. Huang, A. Yang, M. Wu, Y. Liu, X. Dai, and Z. Dong, “Beyond Numerical Time Series: A Unified Benchmark for Multimodal Forecasting with Heterogeneous Context,” arXiv preprint arXiv:2609.15087, 2026. [Online]. Available: <https://arxiv.org/abs/2609.15087>
- [119] T. Zhan, X. Song, T. Guan, S. Pan, and M. Jin, “AION: Next-Generation Tasks and Practical Harness for Time Series,” arXiv preprint arXiv:2605.25045, 2026. [Online]. Available: <https://arxiv.org/abs/2605.25045>
- [120] H. Liu, Y. Zhou, R. Sen, B. A. Prakash, and A. Das, “Rethinking Multimodal Time-Series Forecasting Evaluation,” in ICML 2026, 2026. [Online]. Available: <https://arxiv.org/abs/2607.06973>
- [121] Y. Cai, A. Choudhry, M. Goswami, and A. Dubrawski, “TimeSeriesExam: A time series understanding exam,” in NeurIPS 2024 Workshop, 2024. [Online]. Available: <https://arxiv.org/abs/2410.14752>
- [122] H. Wen, Z. Zhou, Q. Liu, S. Zhong, and Y. Liang, “LiveHouse-TS: An Open-world Living Benchmark for Time Series Foundation Models,” arXiv preprint arXiv:2608.17299, 2026. [Online]. Available: <https://arxiv.org/abs/2608.17299>
- [123] Z. Chen, Y. Wang, J. Wu, Y. Liu, H. Zhang, X. Su, J. Wang, and M. Long, “TimeVista: Exploring and Exploiting Vision-Language Models as Judges for Time Series Forecasting,” arXiv preprint arXiv:2606.16173, 2026. [Online]. Available: <https://arxiv.org/abs/2606.16173>
- [124] S. Nagashima, “Forecast Workflow Bench: Evaluating Language-Model Decisions with Budgeted Forecast Tools,” arXiv preprint arXiv:2609.27385, 2026. [Online]. Available: <https://arxiv.org/abs/2609.27385>
- [125] I. Chakrabarti, Z. Xiong, P. Sharma, and M. Srivastava, “HoliBench: A Cross-Platform Benchmarking and Deployment Toolkit for Foundation Models in CPS-IoT Applications,” arXiv preprint arXiv:2609.12412, 2026. [Online]. Available: <https://arxiv.org/abs/2609.12412>
- [126] T. Ling, C. Qian, and G. Schiele, “Resource-aware Mixed-precision Quantization for Enhancing Deployability of Transformers for Time-series Forecasting on Embedded FPGAs,” arXiv preprint arXiv:2410.03294, 2024. [Online]. Available: <https://arxiv.org/abs/2410.03294>
- [127] K. Yang, W. A. Hanafy, P. Shenoy, and M. Srivastava, “FM-CAC: Carbon-Aware Control for Battery-Buffered Edge AI via Time-Series Foundation Models,” arXiv preprint arXiv:2604.16448, 2026. [Online]. Available: <https://arxiv.org/abs/2604.16448>
- [128] H. Wang, R. Xu, G. Kementzidis, K. Cho, S. R. Villarrreal, and Y. Deng, “Towards Principled Test-Time Adaptation for Time Series Forecasting,” arXiv preprint arXiv:2605.17250, 2026. [Online]. Available: <https://arxiv.org/abs/2605.17250>
- [129] Y. Cheng, Z. Wang, T. Xiao, Q. Zhong, J. Suo, and K. He, “CausalTime: Realistically Generated Time-series for Benchmarking of Causal Discovery,” in NeurIPS 2023, 2023. [Online]. Available: <https://arxiv.org/abs/2310.01753>
- [130] A. Garza, R. Rosillo, R. Mendoza-Smith, D. Salinas, A. R. Williams, A. Ashok, M. Goswami, and J. M. Juárez, “Impermanant: A Live Benchmark for Temporal Generalization in Time Series Forecasting,” arXiv preprint arXiv:2603.08707, 2026. [Online]. Available: <https://arxiv.org/abs/2603.08707>
- [131] H. Li, S. Xie, L. Shen, Z. Li, M. Chen, X. Zhang, H. Fu, J. Sun, X. Ren, and C. Liu, “TSFMAudit: Data Contamination

- Auditing in Forecasting Time Series Foundation Models,” arXiv preprint arXiv:2605.26161, 2026. [Online]. Available: <https://arxiv.org/abs/2605.26161>
- [132] M. N. Moghadasi and F. Ghaderi, “A Later Test Set Is Not a New Domain: Pretraining Familiarity Survives a Contamination-Free Hold-Out,” arXiv preprint arXiv:2609.10357, 2026. [Online]. Available: <https://arxiv.org/abs/2609.10357>
 - [133] Z. Pan and A. A. Ezzat, “Evaluating Time Series Foundation Models for Electricity Price Forecasting: Contamination Risk, Distributional Shifts, and Covariate Dependence,” arXiv preprint arXiv:2607.02623, 2026. [Online]. Available: <https://arxiv.org/abs/2607.02623>
 - [134] L. Ye, H. Wang, F. Ahmed, R. Zhang, Y. He, Z. Lin, Y. Yin, J. Peng, M. Galarzyk, S. Chava, and C. Zhang, “Forecast-Dojo: Replayable Environments for Benchmarking and Training LLM Forecasting Agents,” arXiv preprint arXiv:2609.28876, 2026. [Online]. Available: <https://arxiv.org/abs/2609.28876>
 - [135] S. Achour, Y. Bouher, D. Nguyen, and N. Chesneau, “Foundation models for time series forecasting: Application in conformal prediction,” arXiv preprint arXiv:2507.08858, 2025. [Online]. Available: <https://arxiv.org/abs/2507.08858>
 - [136] N. M. Gil, F. O’Donncha, W. M. Gifford, N. Zhou, D. C. Patel, and R. Vaculin, “Adaptive Conformal Anomaly Detection with Time Series Foundation Models for Signal Monitoring,” in ICLR 2026, 2026. [Online]. Available: <https://arxiv.org/abs/2604.20122>
 - [137] M. Heurich, M. Granz, and T. Landgraf, “RareCP: Regime-Aware Retrieval for Efficient Conformal Prediction,” arXiv preprint arXiv:2605.08857, 2026. [Online]. Available: <https://arxiv.org/abs/2605.08857>